

Begin your Deep Learning Journey with Art



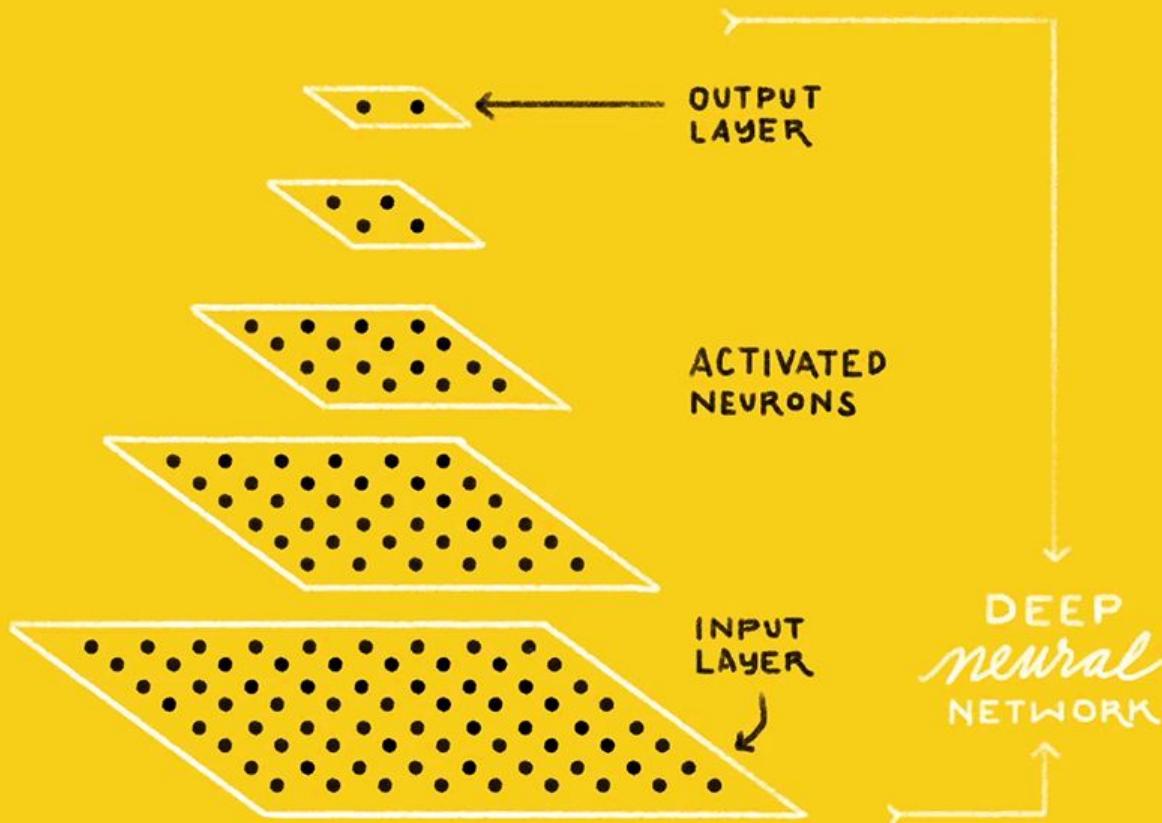
Josh Gordon
@random_forests

Messe Karlsruhe

IS THIS A
CAT or DOG?



CAT DOG



Deep Learning

Current state of the art in many tasks:

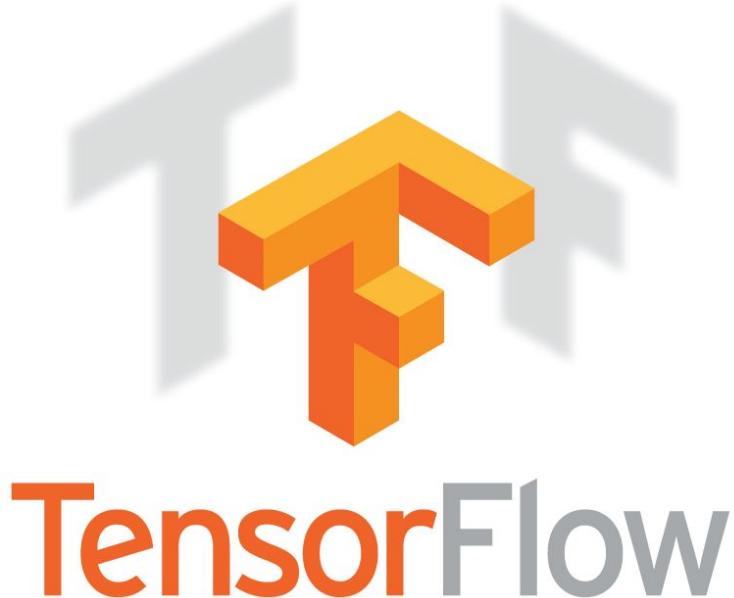
Image: *classification, captioning, colorization*

Text: *translation, parsing, summarization*

Speech: *recognition, generation*

Games: *AlphaGo, Atari, Mario Kart*

Many more.



- **Fast, flexible, and scalable open-source machine learning library**
- **For research and production**
- **Runs on CPU, GPU, Android, iOS, Raspberry PI, Datacenters**
- **Apache 2.0 license**

<https://research.googleblog.com/2016/11/celebrating-tensorflows-first-year.html>

Reproducible Research

Music <https://magenta.tensorflow.org/2016/11/09/tuning-recurrent-networks-with-reinforcement-learning/>

WaveNet <https://deepmind.com/blog/wavenet-generative-model-raw-audio/>

Translation <https://research.googleblog.com/2016/09/a-neural-network-for-machine.html>

Summarization <https://research.googleblog.com/2016/08/text-summarization-with-tensorflow.html>

Show and Tell <https://research.googleblog.com/2016/09/show-and-tell-image-captioning-open.html>

Inception <https://research.googleblog.com/2016/08/improving-inception-and-image.html>

Parsey McParseface <https://research.googleblog.com/2016/05/announcing-syntaxnet-worlds-most.html>

Many more: <https://github.com/tensorflow/models/>

Images

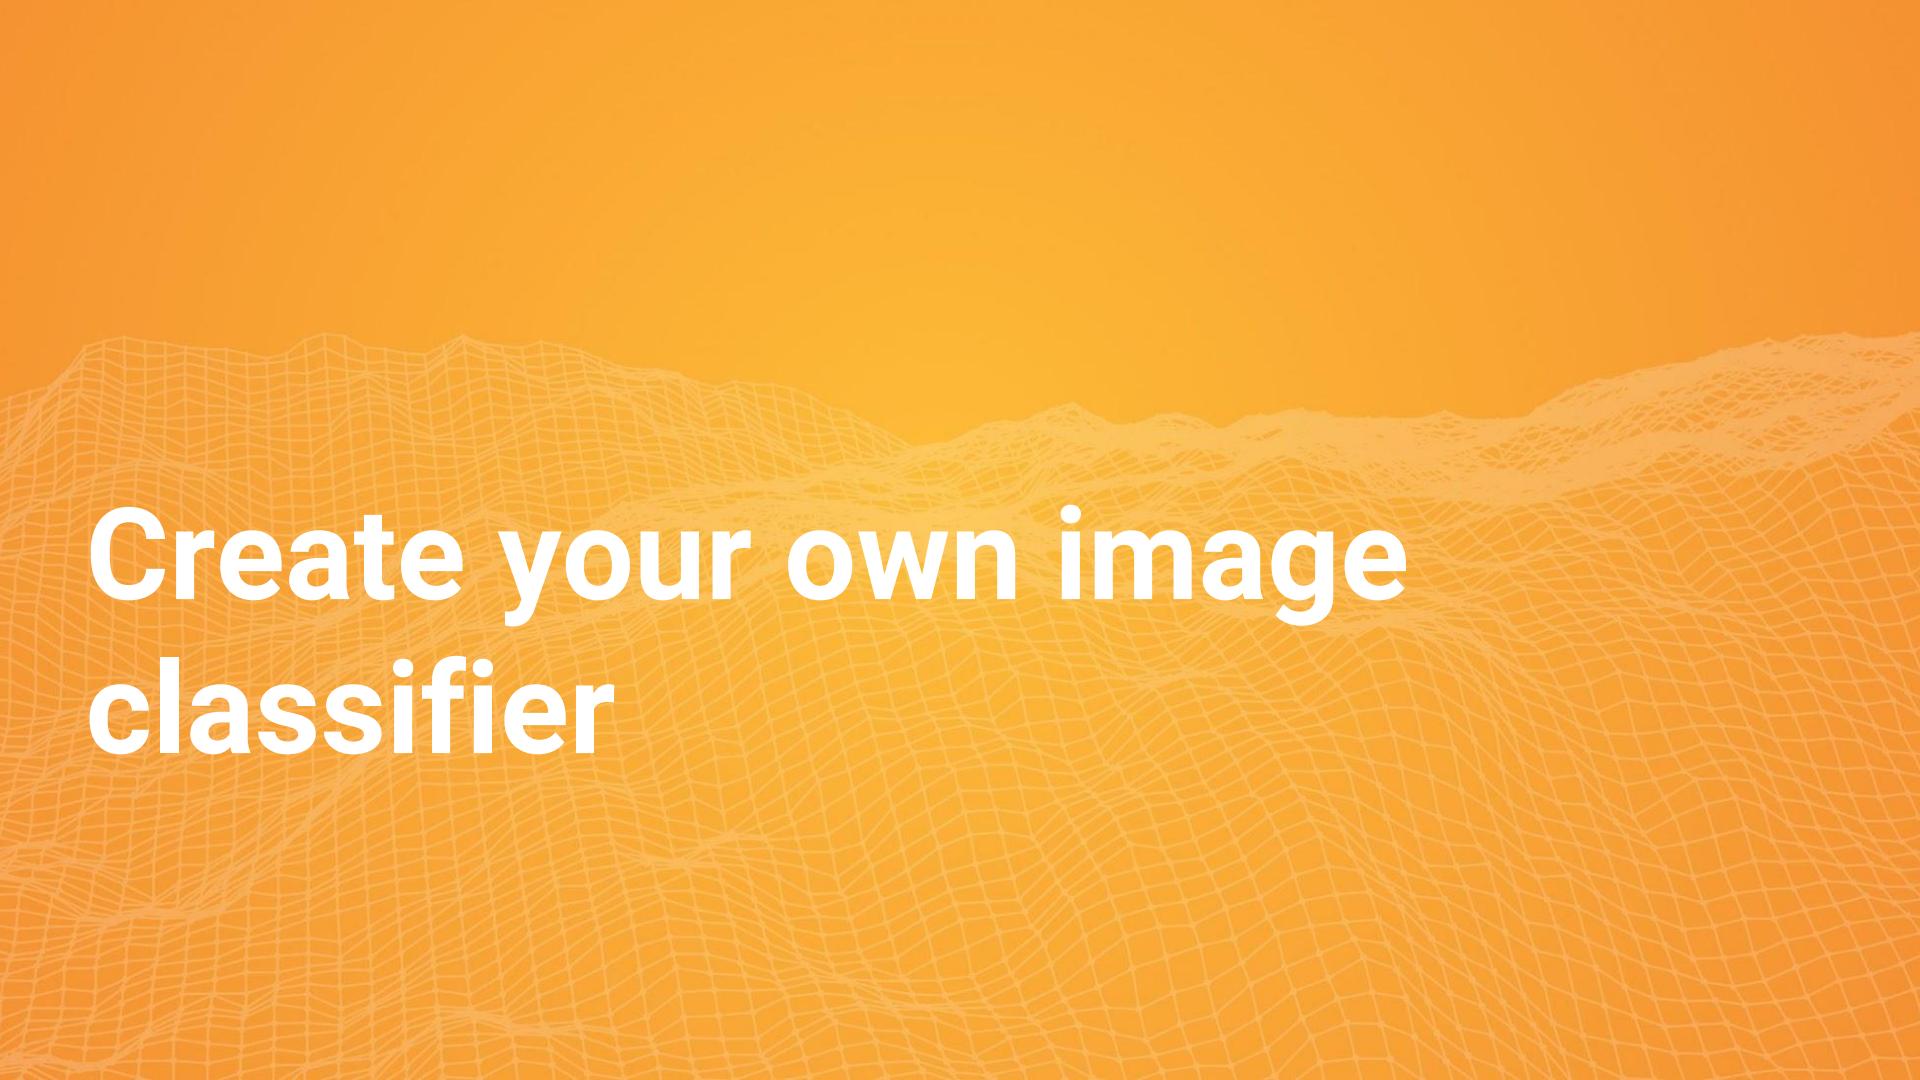


Inception



An [Alaskan Malamute](#) ([left](#)) and a [Siberian Husky](#) ([right](#)). Images from Wikipedia.

<https://research.googleblog.com/2016/08/improving-inception-and-image.html>



Create your own image classifier

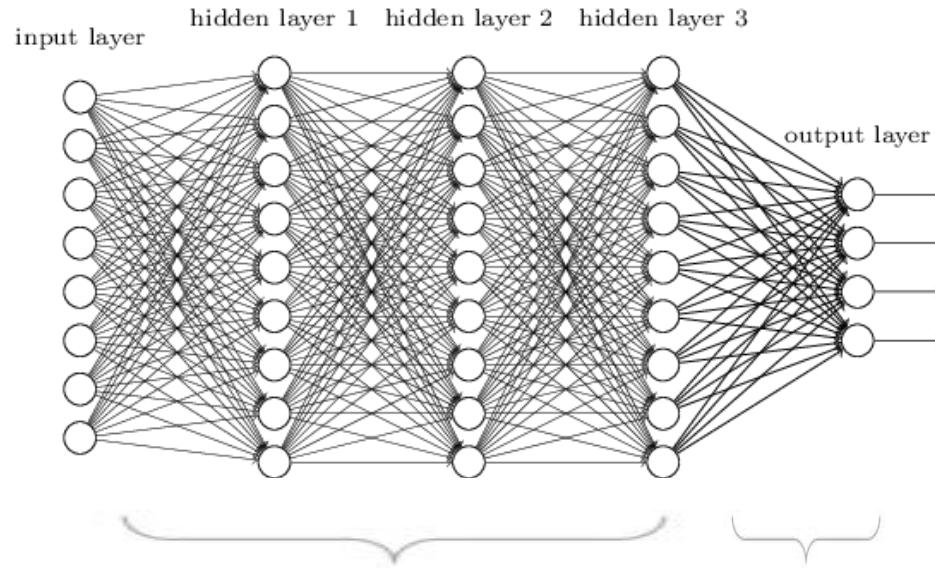
goo.gl/xGsB9d



Claude Monet - Bouquet of Sunflowers
Images from the Metropolitan Museum of Art (with permission)

For more: [@random_forests](https://twitter.com/random_forests)

Transfer Learning

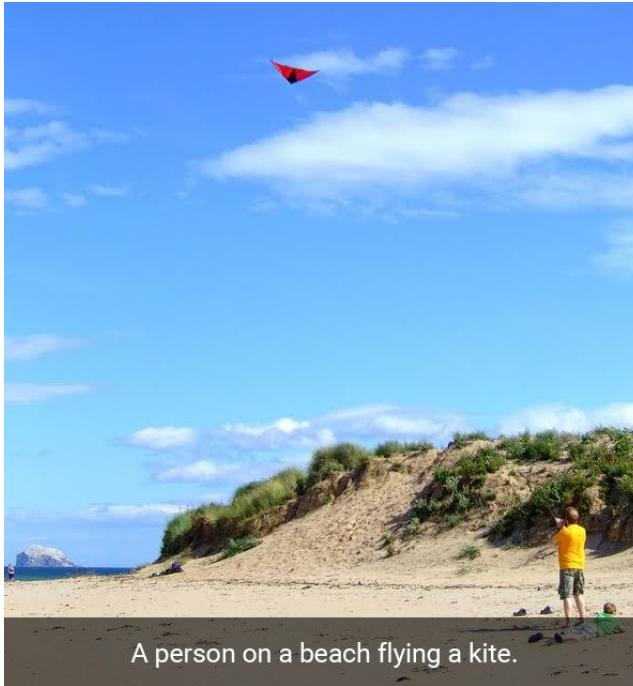


Most of the network is unchanged Only the last layer of weights
and the outputs are updated

TensorFlow for Poets

Codelab goo.gl/xGsB9d
Video goo.gl/KewA03

Show and Tell

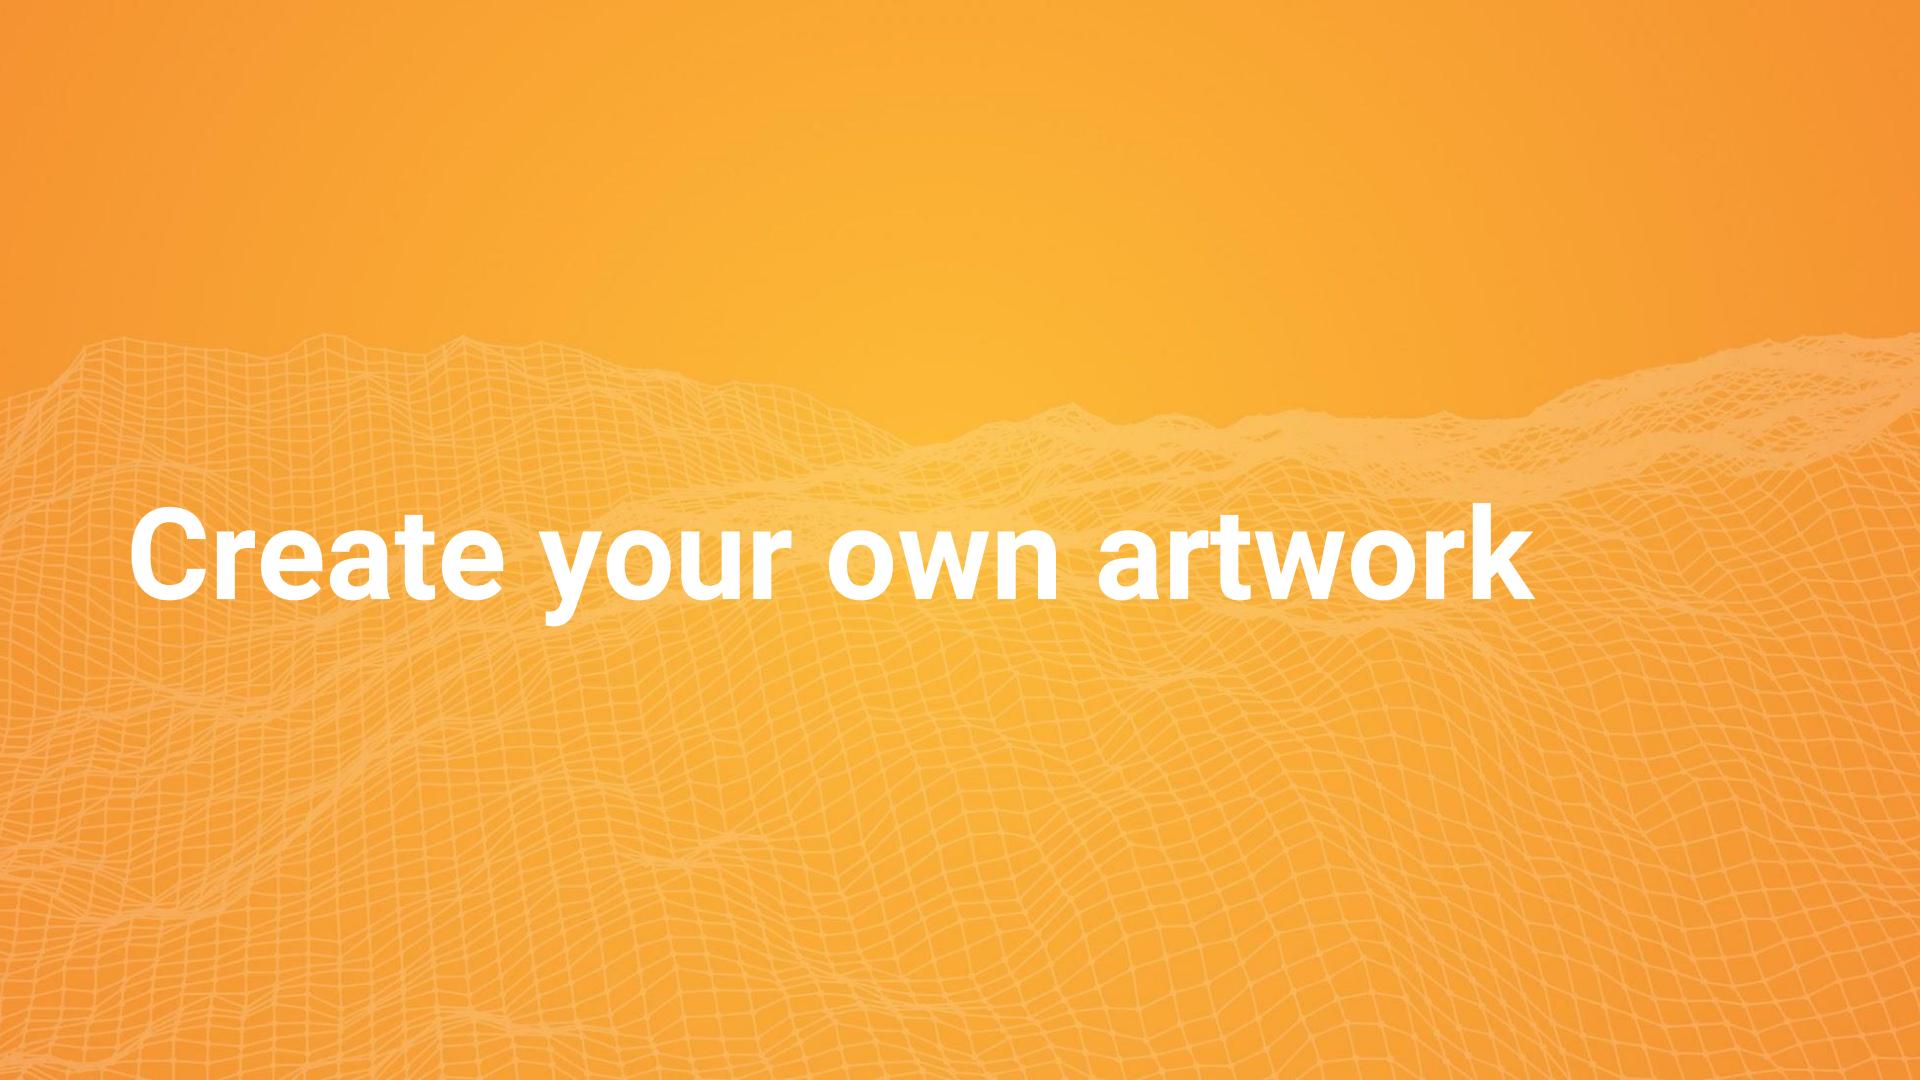


<https://research.googleblog.com/2016/09/show-and-tell-image-captioning-open.html>

Pretrained model - goo.gl/ExBjhS

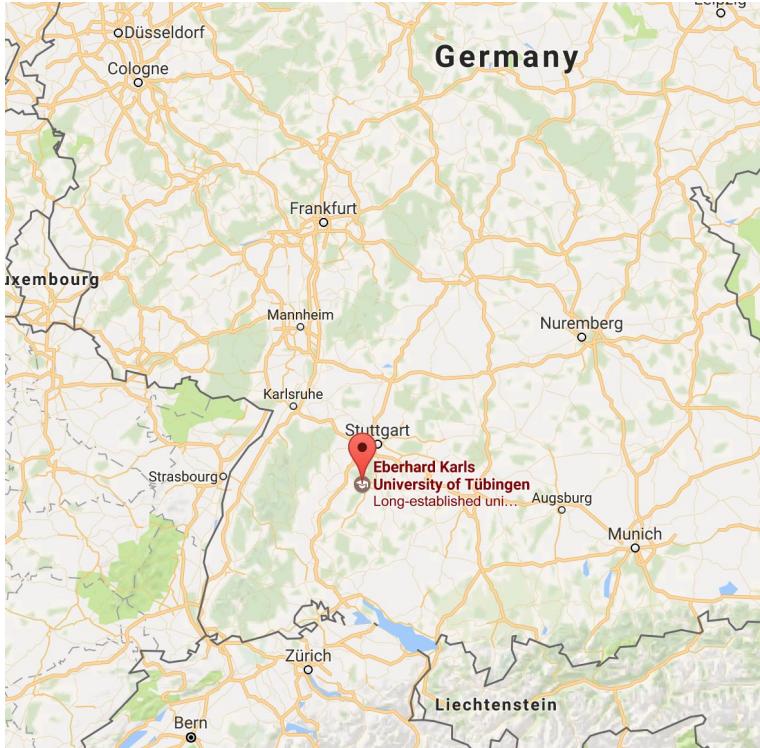
A dog with a frisbee
in its mouth



The background features a solid orange gradient at the top transitioning into a white wireframe mesh pattern that covers the entire frame. The wireframe consists of a grid of thin, light-colored lines forming a series of undulating hills or waves.

Create your own artwork

A Neural Algorithm of Artistic Style



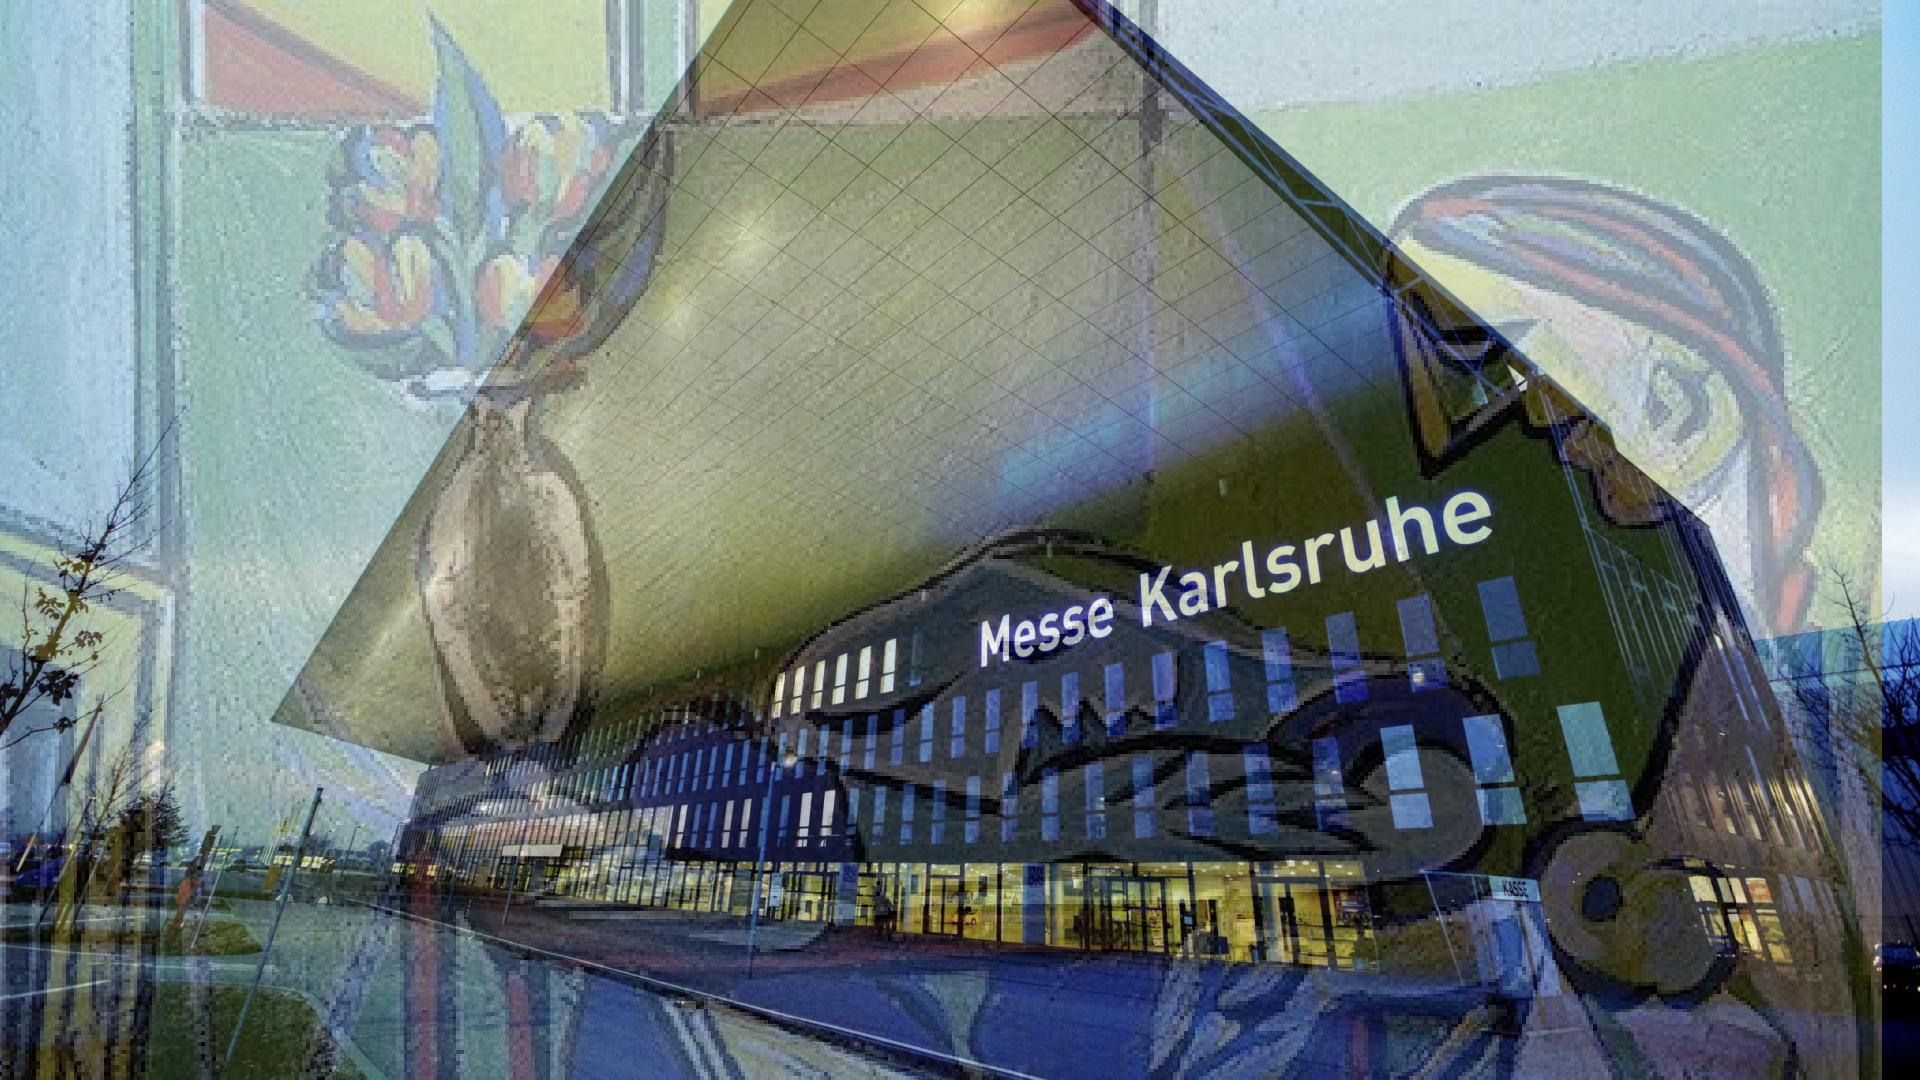
<https://arxiv.org/pdf/1508.06576v2.pdf>

A wide-angle photograph of the Messe Karlsruhe building at dusk. The building features a large, angular, light-colored roof and a dark facade with many illuminated windows. The words "Messe Karlsruhe" are prominently displayed in white, illuminated letters along the side of the building. The sky is a clear blue.

Messe Karlsruhe

KASSE





Messe Karlsruhe

A colorful, abstract painting of the Messe Karlsruhe building. The building features a large, multi-story facade with numerous windows. A prominent feature is a large, triangular roof section with a grid-like pattern of colored tiles in shades of blue, purple, and yellow. The words "Messe Karlsruhe" are written in white, stylized letters across the upper part of the building's facade. The overall style is vibrant and textured, with visible brushstrokes and a mix of warm and cool colors.

Messe Karlsruhe

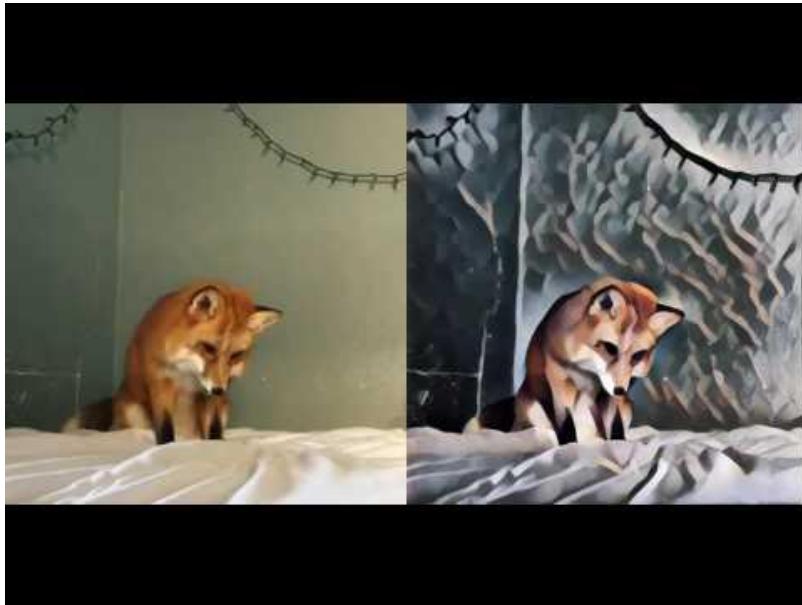
富嶽三十六景 神奈川沖浪裏

一七九九年





Messe Karlsruhe

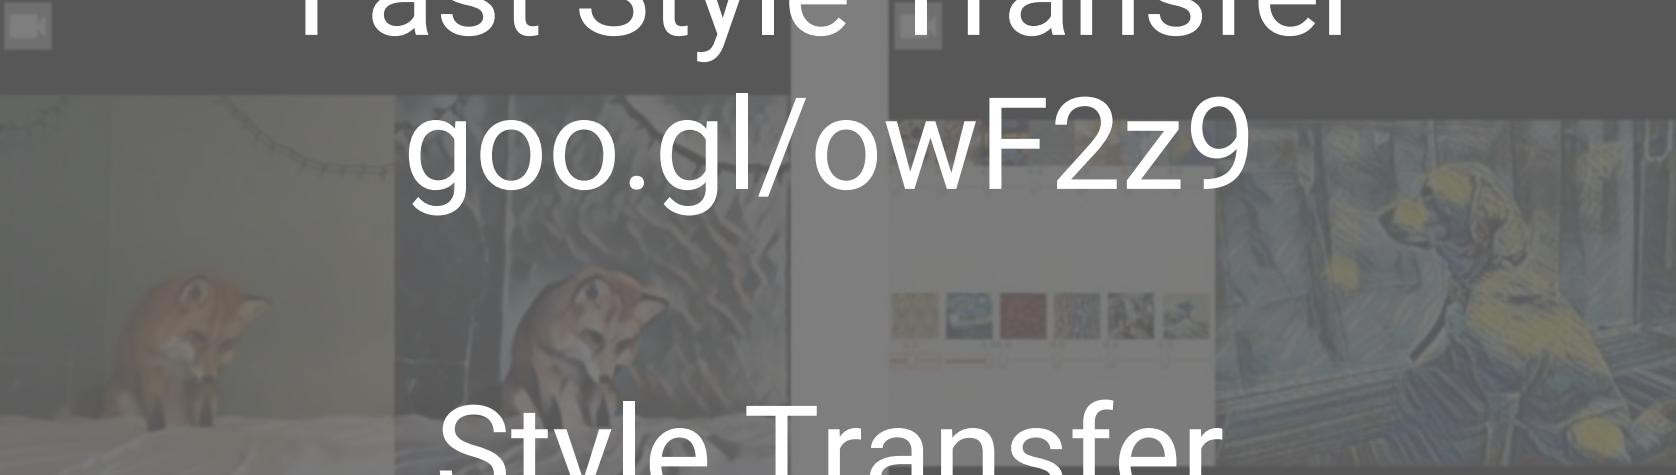


<https://github.com/lengstrom/fast-style-transfer/>

<https://research.googleblog.com/2016/10/supercharging-style-transfer.html>

Fast Style Transfer

goo.gl/owF2z9

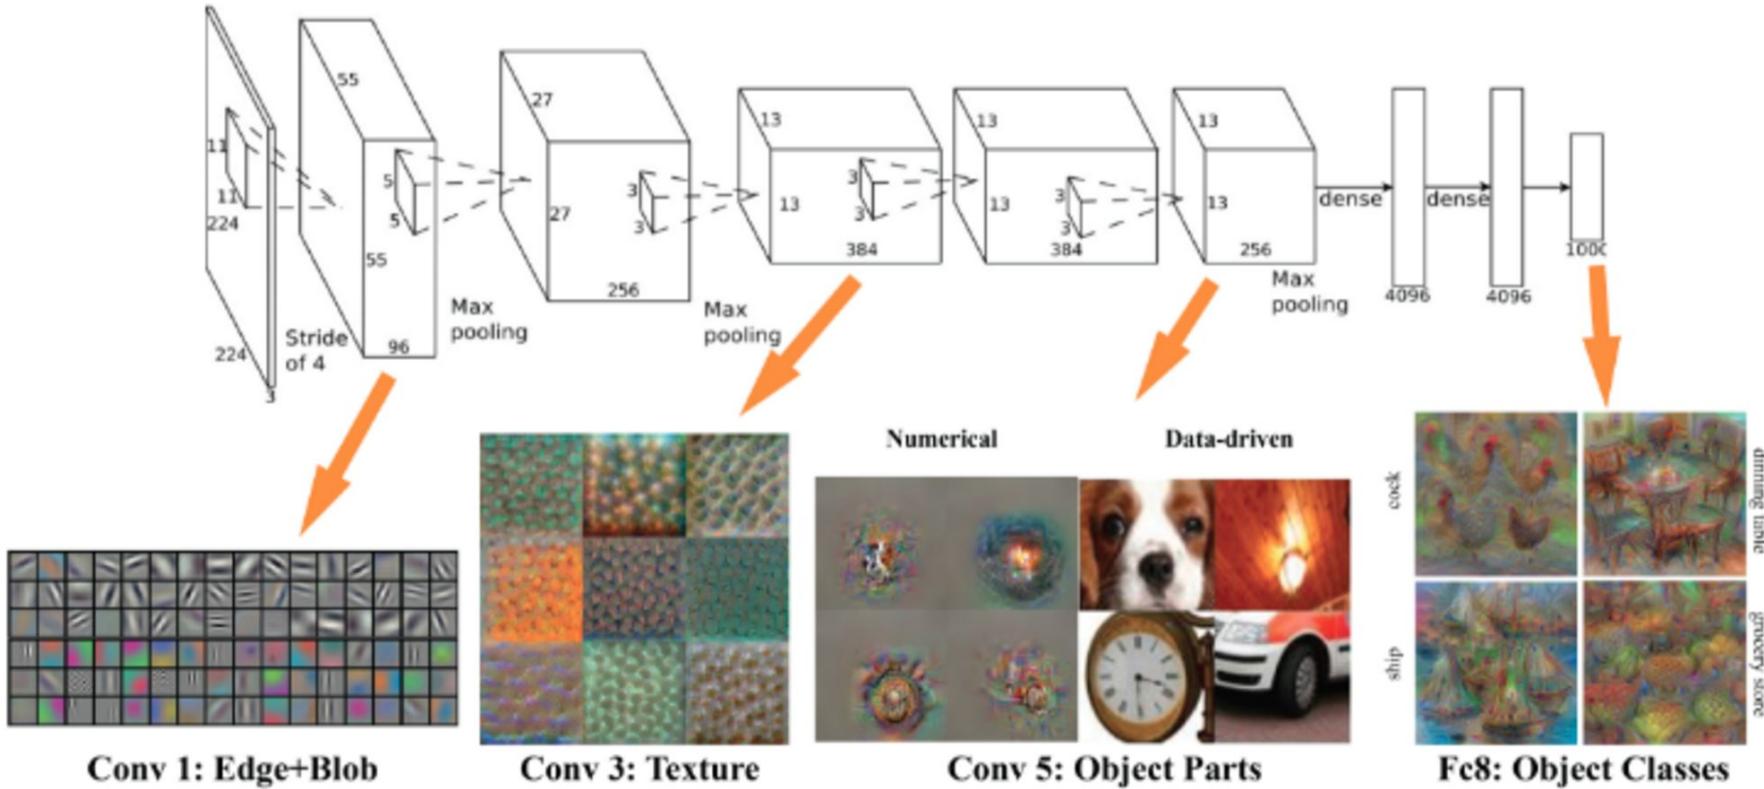


Style Transfer

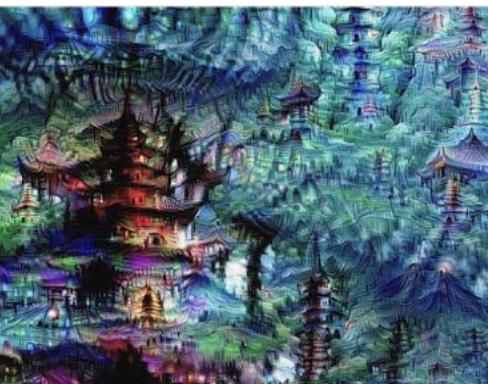
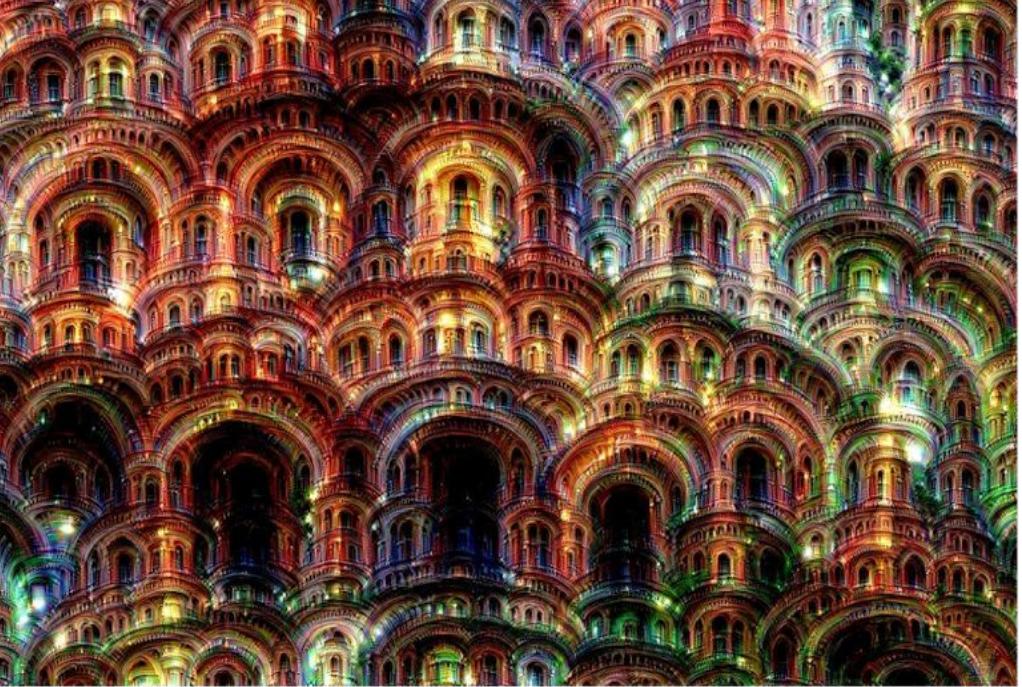
goo.gl/M5hiYY

<https://github.com/lengstrom/fast-style-transfer/>

<https://research.googleblog.com/2016/10/supercharging-style-transfer.html>



From: [A Matlab Plugin to Visualize Neurons from Deep Models](#), Donglai Wei et. al.



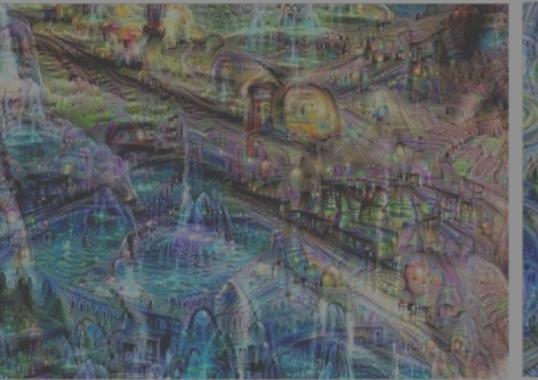
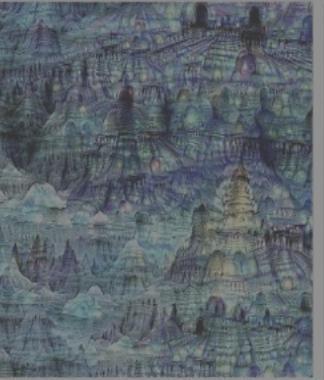


<https://www.youtube.com/watch?v=DgPaCWJL7XI>

<https://research.googleblog.com/2015/06/inceptionism-going-deeper-into-neural.html>



goo.gl/1kBXyO



ColorNet

Grayscale

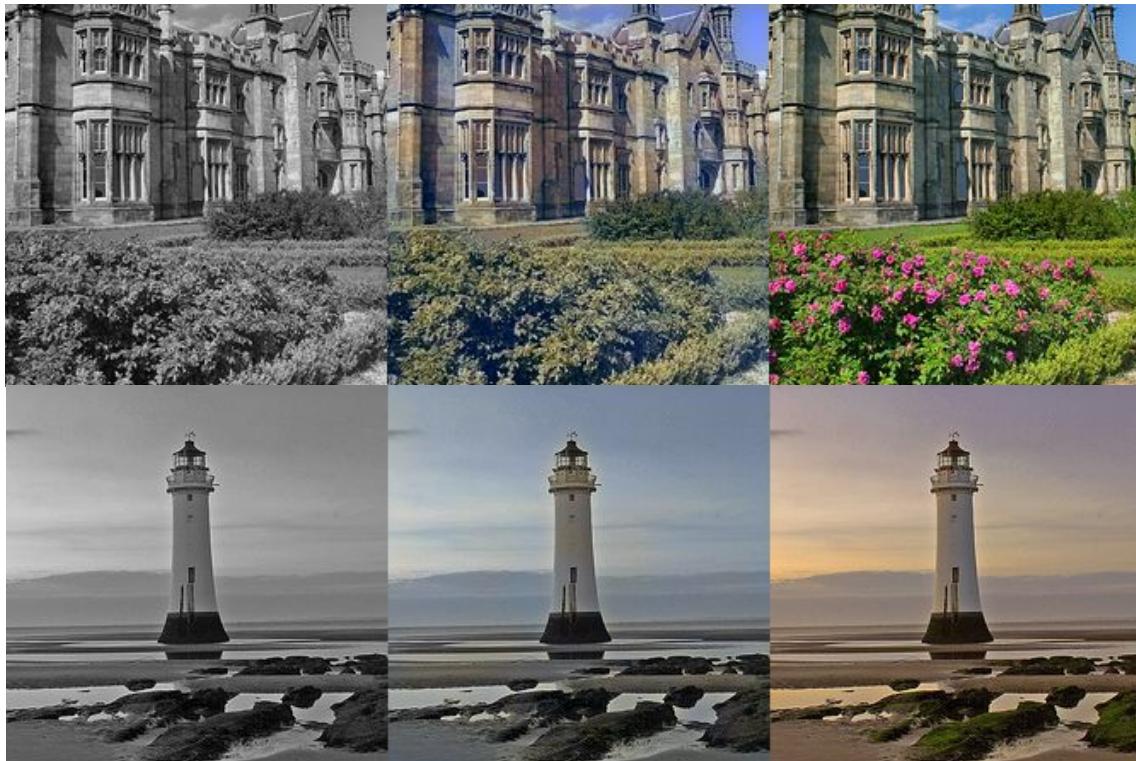
Prediction

Ground Truth

灰度图

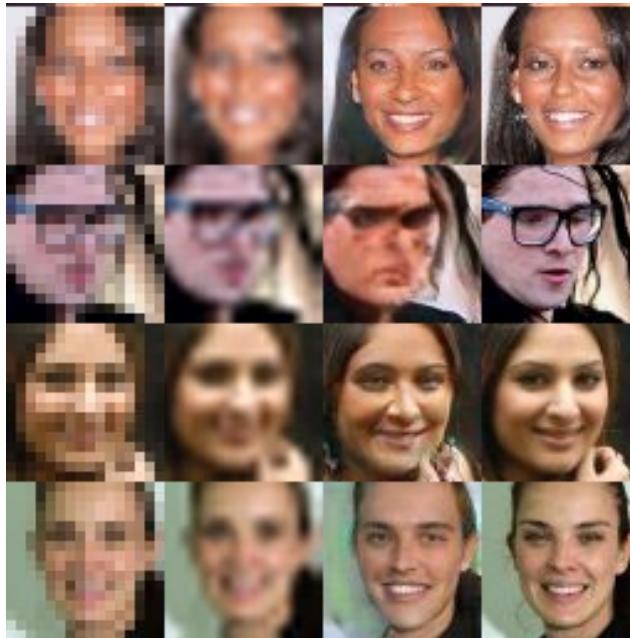
预测

地面实况



<https://github.com/pavelgonchar/colornet>

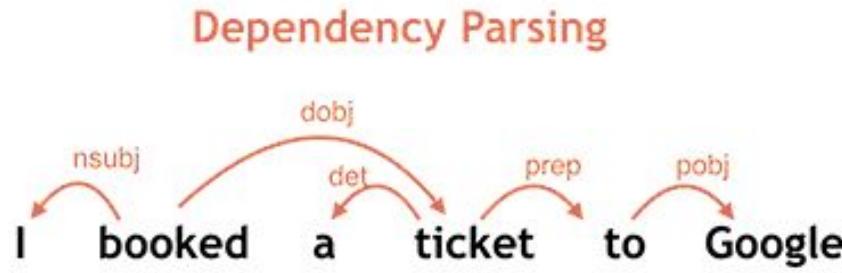
Image Super-Resolution



<https://github.com/david-gpu/srez>

Text

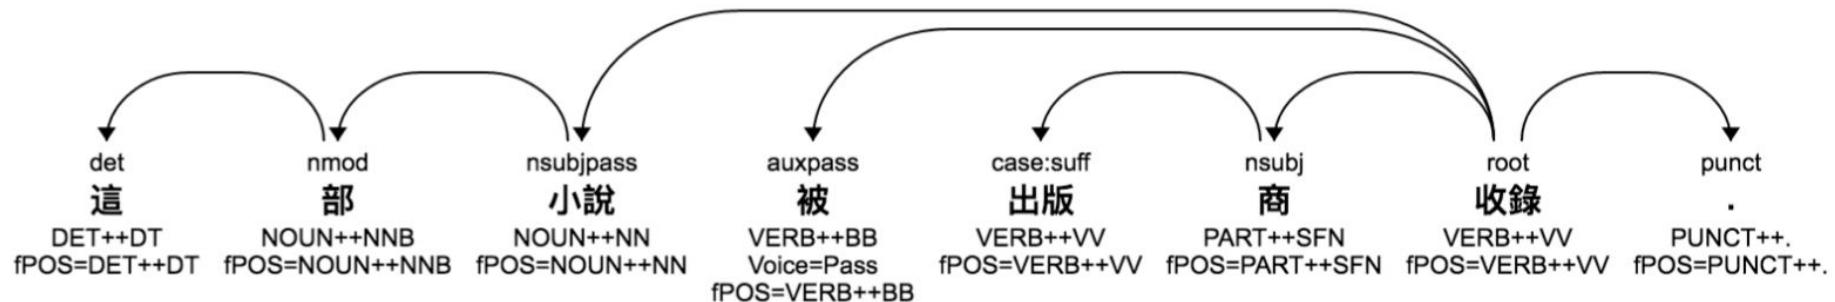
Parsey McParseface



我 订了 一张 票 去 Google

Meet Parsey's Cousins: Syntax for 40 languages

這部小說被出版商收錄。



Summarization

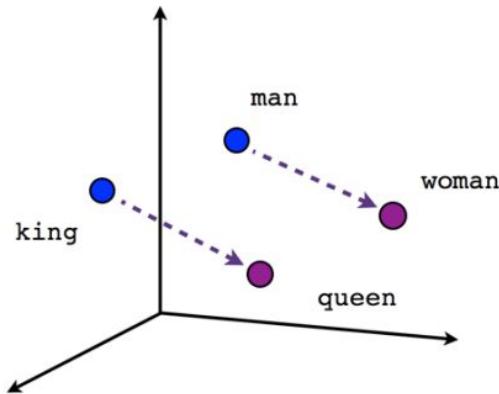
Original text

- *Alice and Bob took the train to visit the zoo. They saw a **baby giraffe, a lion, and a flock of colorful tropical birds.***

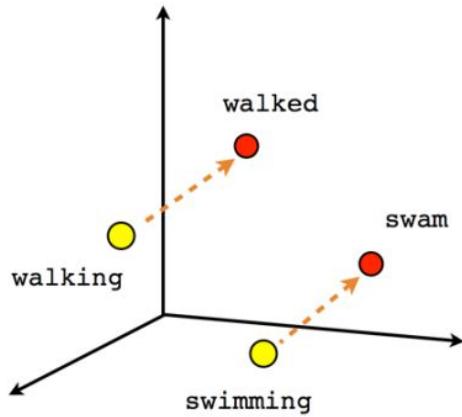
Abstractive summary

- *Alice and Bob visited the zoo and saw **animals and birds.***

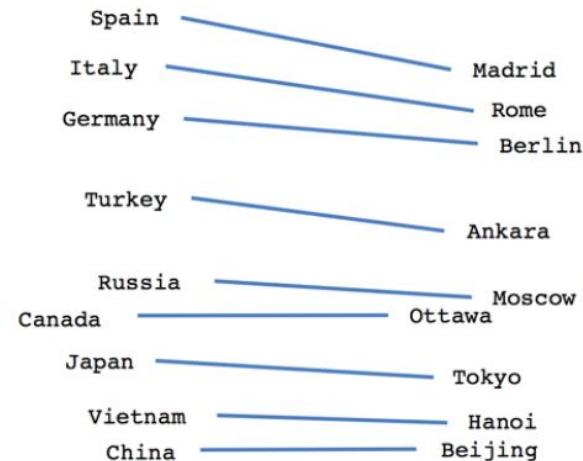
Word Embeddings



Male-Female

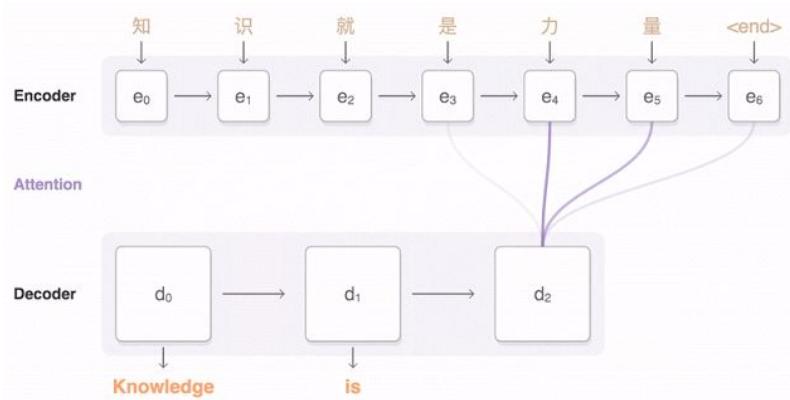


Verb tense



Country-Capital

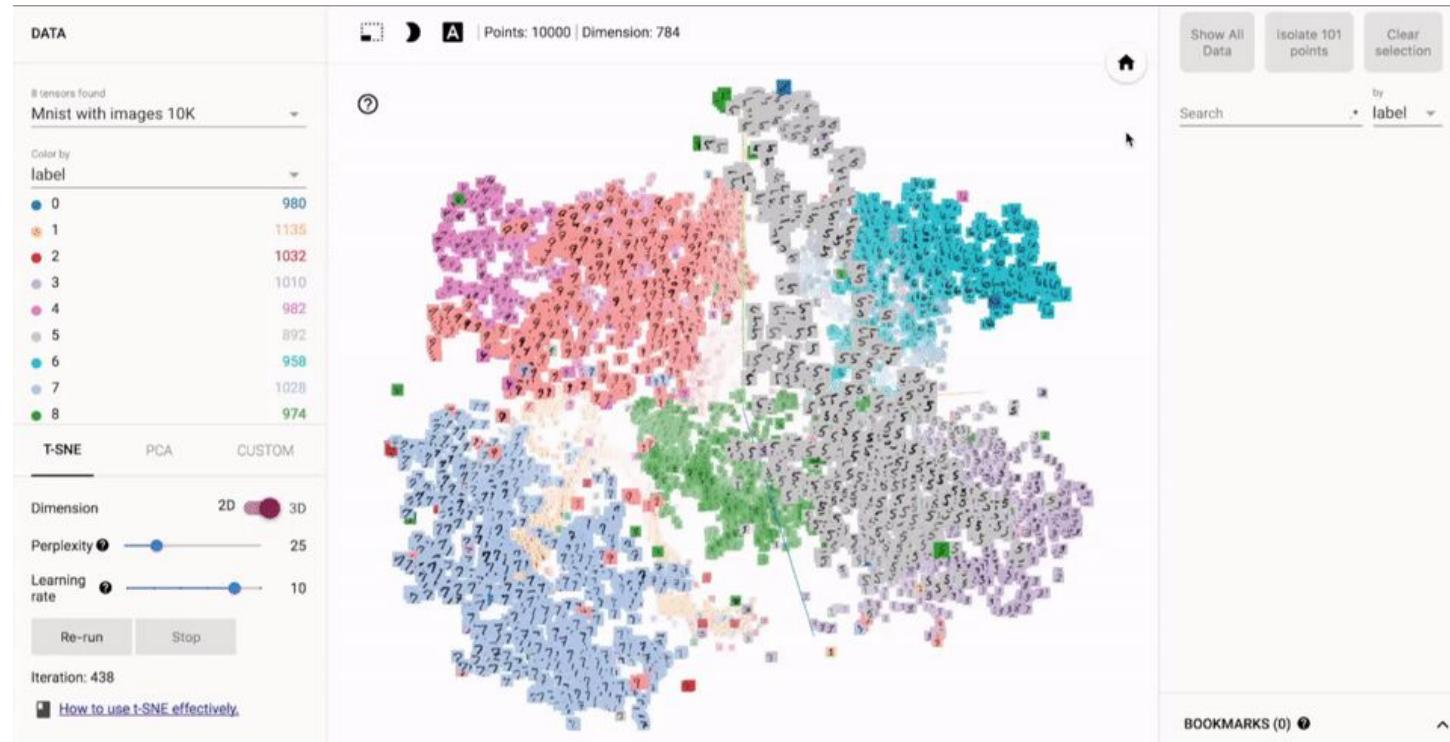
Translation



Chinese->English	100年前，预测引力波的爱因斯坦或许都无法想象人类可以直接观测到引力波。	100 years ago, the prediction of Einstein's gravitational waves probably can not imagine humans can directly observe gravitational waves.	100 years ago, Einstein predicted gravitational waves may not be able to imagine humans can directly observe the gravitational waves.	100 years ago, Einstein who predicted gravitational waves may not be able to imagine that humans can directly observe the gravitational waves.
------------------	--------------------------------------	---	---	--

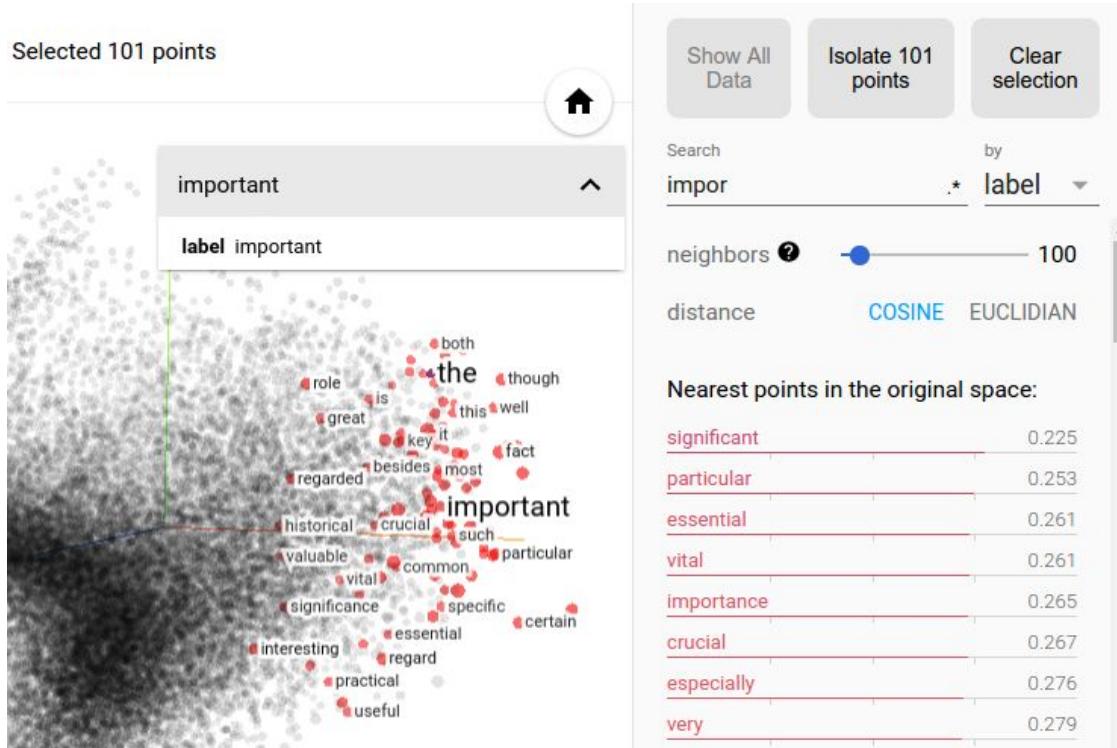
<https://research.googleblog.com/2016/09/a-neural-network-for-machine.html>

TensorBoard



https://www.tensorflow.org/versions/r0.12/how_tos/embedding_viz/index.html

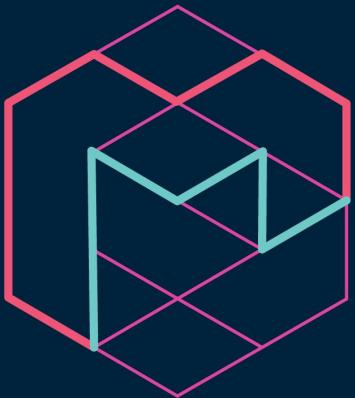
TensorBoard



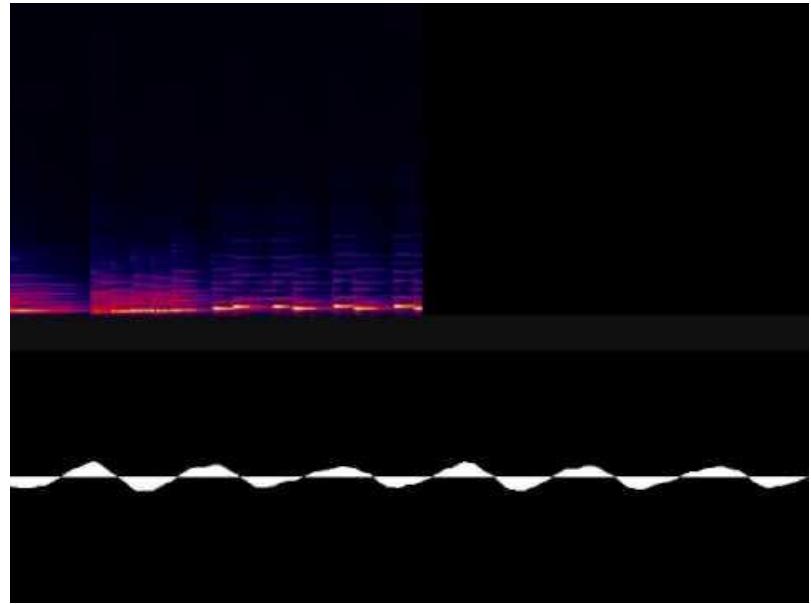
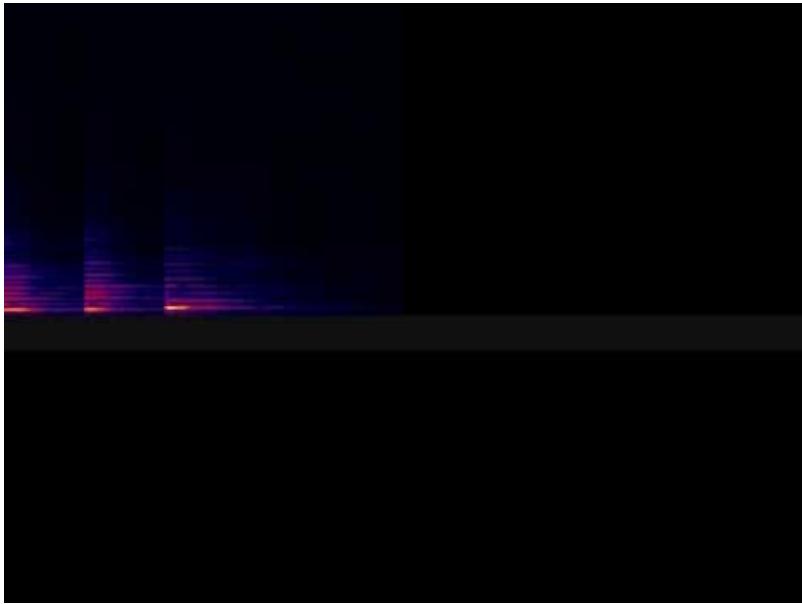
https://www.tensorflow.org/versions/r0.12/how_tos/embedding_viz/index.html

Sound





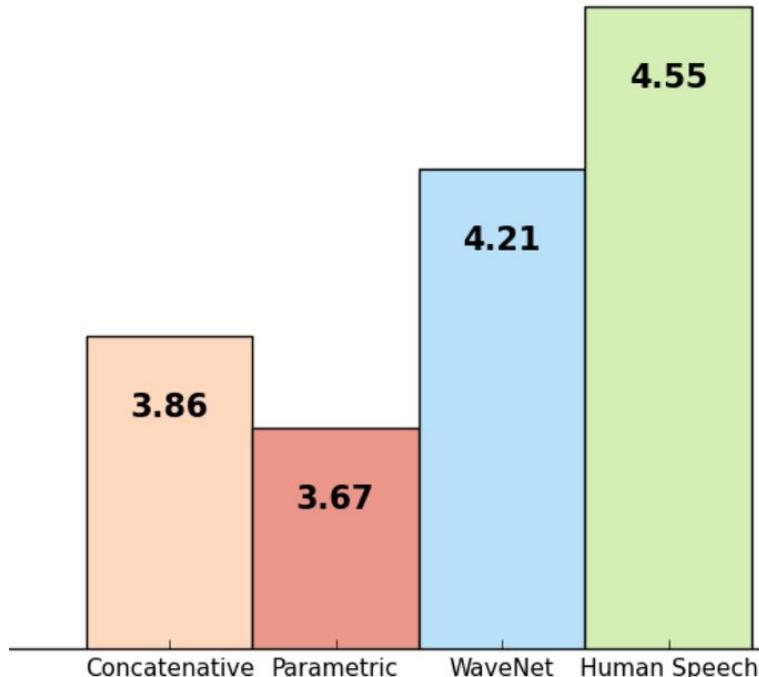
MAGENTA



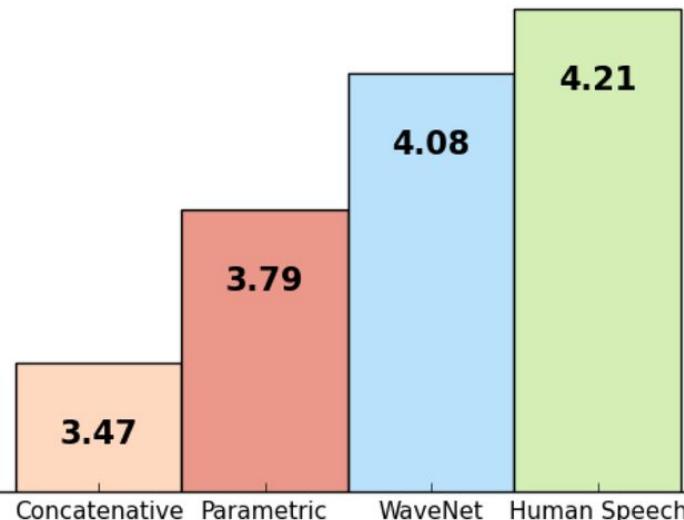
<https://magenta.tensorflow.org/2016/11/09/tuning-recurrent-networks-with-reinforcement-learning/>

WaveNet

US English



Mandarin Chinese



<https://deepmind.com/blog/wavenet-generative-model-raw-audio/>

Generating Speech and Music

[English – Parametric 英文 – 参数化](#)

[English – Concatenative 英文 - 拼接](#)

[English – WaveNet 英文 - WaveNet](#)

[Mandarin – Parametric 汉语 - 参数化](#)

[Mandarin – Concatenative 汉语 – 拼接](#)

[Mandarin – WaveNet 汉语 - WaveNet](#)

[Bonus - on hold 福利时间 – 听一段铃声](#)

<https://deepmind.com/blog/wavenet-generative-model-raw-audio/>

<https://magenta.tensorflow.org/2016/11/09/tuning-recurrent-networks-with-reinforcement-learning/>

From the community

A TensorFlow implementation of DeepMind's WaveNet paper

<https://github.com/ibab/tensorflow-wavenet>

Speech-to-Text-WaveNet : End-to-end sentence level English speech
recognition based on DeepMind's WaveNet and TensorFlow

<https://github.com/buriburisuri/speech-to-text-wavenet>

Games





BBC News Sport Weather iPlayer TV Radio More Search

NEWS

Find local news

Home UK World Business Politics Tech Science Health Education Entertainment & Arts More

Technology

Google achieves AI 'breakthrough' at Go

An artificial intelligence program developed by Google beats Europe's top player at the ancient Chinese game of Go, about a decade earlier than expected.

© 27 January 2016 | Technology

How did they do it?
What is the game Go?

Facebook trains AI to beat humans at Go



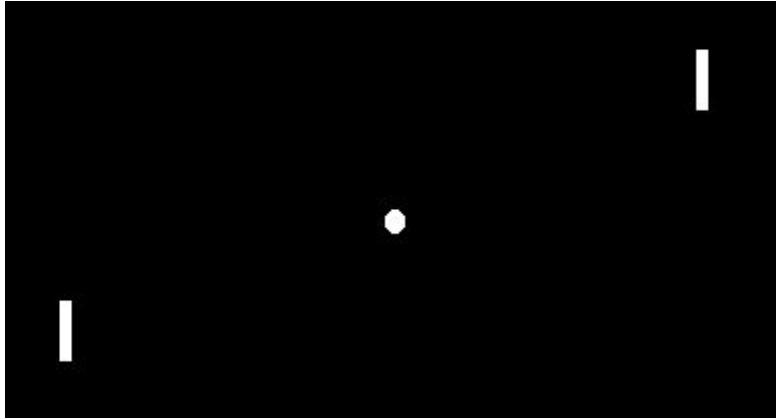
Google's AI just cracked the game that supposedly no computer could beat

By Mike Murphy | January 27, 2016



Going up. (Reuters/Kiyoshi Ota)

Computers have slowly started to encroach on activities we previously believed only the brilliantly sophisticated human brain could handle. IBM's Deep Blue supercomputer beat Grand Master Garry Kasparov at chess in 1997, and in 2011 IBM's Watson beat former human winners at the quiz game *Jeopardy*. But the ancient board game Go has long been one of the major goals of artificial intelligence research. It's understood to be one of the most difficult games for computers to handle due to the sheer number of possible moves a player can make at any given point. Until now, that is.



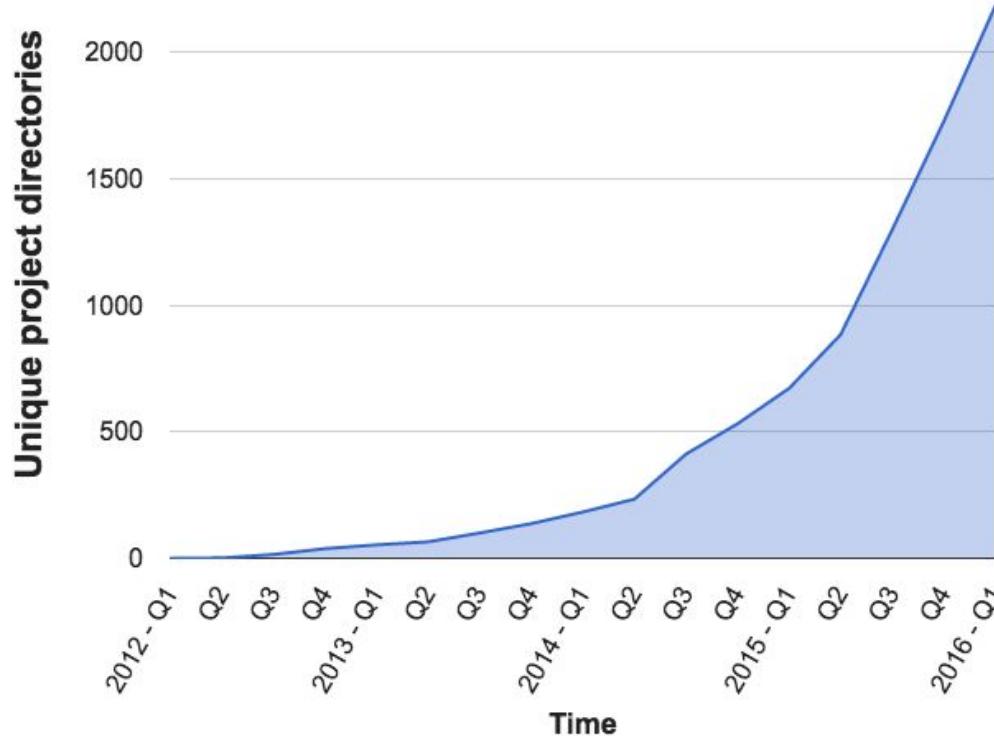


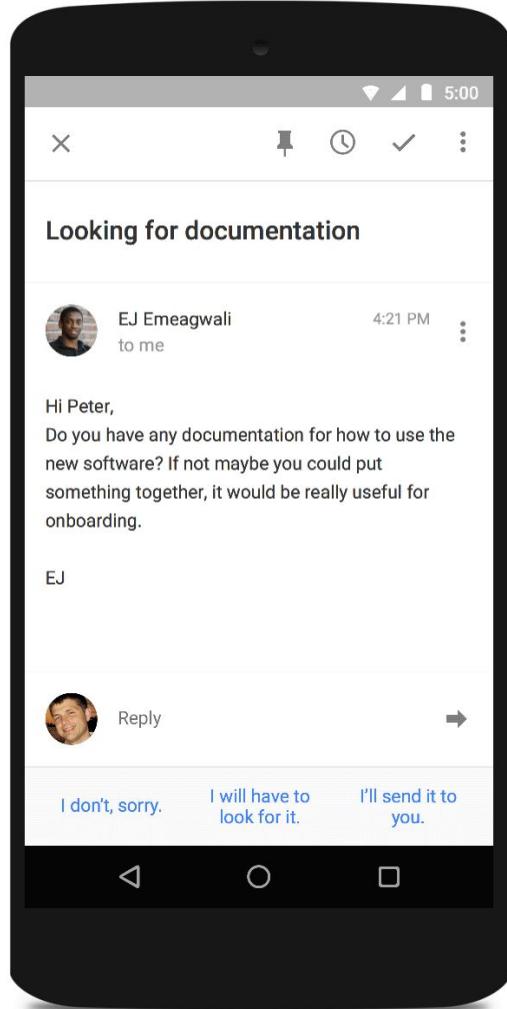
<http://kevinhughes.ca/blog/tensor-kart>

Deep Learning at Google

Growing Use of Deep Learning at Google

of directories containing model description files





Smart reply in Inbox by Gmail

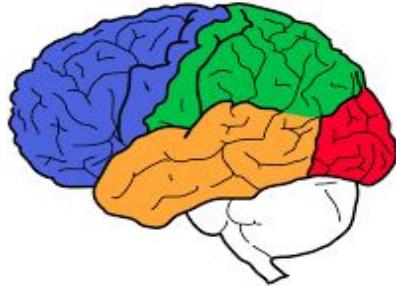
10%

of all responses
sent on mobile

In conclusion

“Universal” Machine Learning

*Speech
Text
Search
Queries
Images
Videos
Labels
Entities
Words
Audio
Features*



*Speech
Text
Search
Queries
Images
Videos
Labels
Entities
Words
Audio
Features*

Detection of Diabetic Eye Disease



<https://research.googleblog.com/2016/11/deep-learning-for-detection-of-diabetic.html>

Next steps

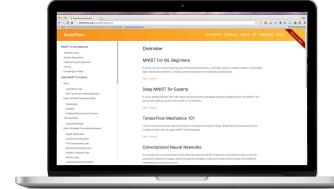
Tutorials and code
tensorflow.org

Totally new to ML?
Recipes goo.gl/KewA03

Intro to Deep Learning with TensorFlow
Udacity class goo.gl/iHssl

Stanford's CS231n
cs231n.github.io

Udacity's Machine Learning Nanodegree
goo.gl/ODpXj4



These slides + exercises
goo.gl/nrdsxM

TensorFlow Playground
goo.gl/mXhncM

Chris Olah's blog
colah.github.io

Michael Nielsen's book
neuralnetworksanddeeplearning.com

Thank you and have fun!



Josh Gordon
[@random_forests](https://twitter.com/random_forests)

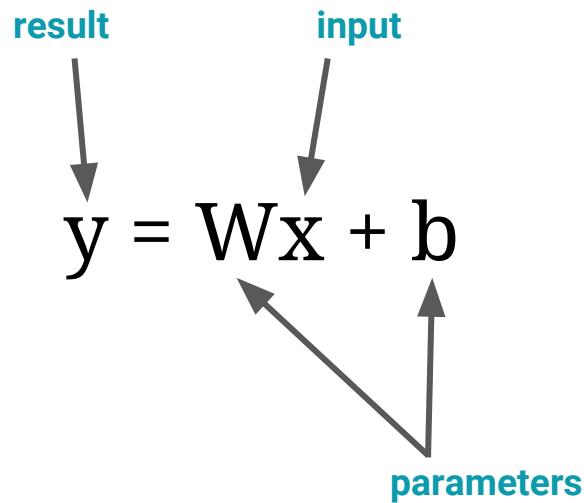
Extras

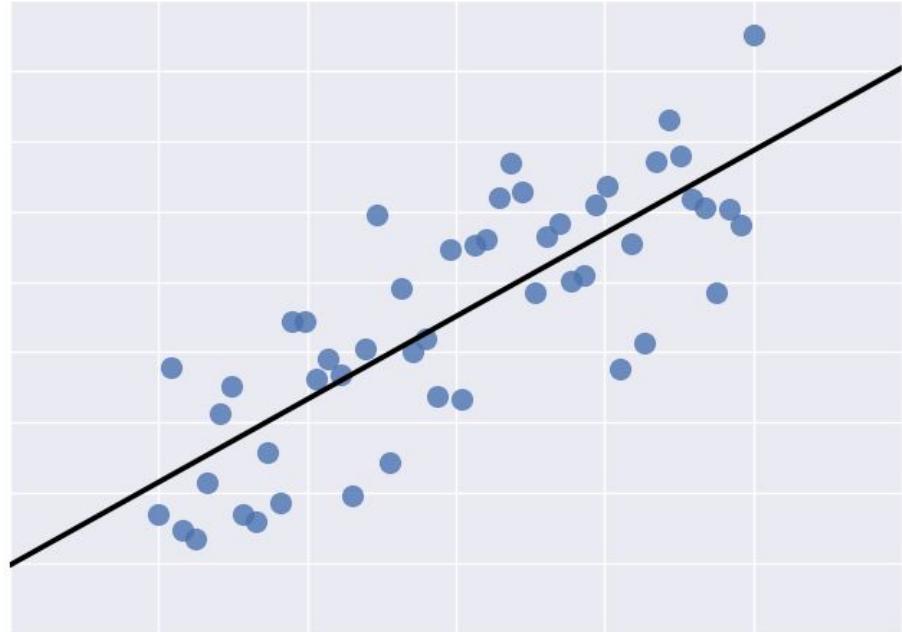


Linear Regression

$$y = Wx + b$$

result input
parameters





Let's code this up

1. **Inference** (“given x , this is the code to predicts y ”)
2. **Loss** (“quantify how bad our prediction was”)
3. **Optimize** (“update variables to improve the prediction”)

Inference: putting it all together

```
import tensorflow as tf
x = tf.placeholder(shape=[None],
                    dtype=tf.float32,
                    name='x')
W = tf.Variable(tf.random_normal([1], name="W"))
b = tf.Variable(tf.random_normal([1], name="b"))
y = W * x + b

with tf.Session() as sess:
    sess.run(tf.initialize_all_variables())
    print(sess.run(y, feed_dict={x: x_in}))
```

Build the graph

Prepare execution env

Initialize variables

Run the computation

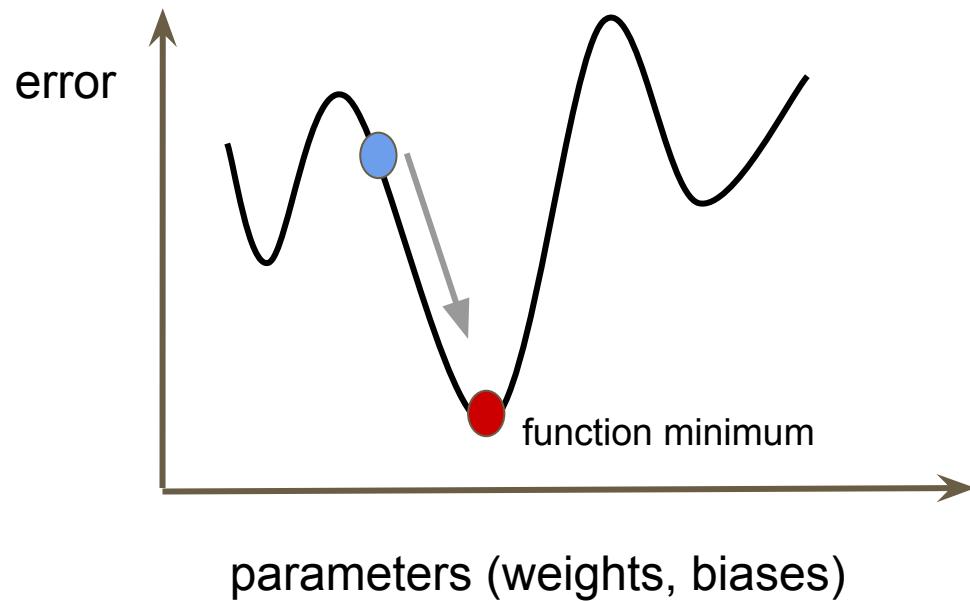
Define a Loss

Given y , y_{train} compute a loss, for instance:

$$loss = (y - y_{train})^2$$

```
# create an operation that calculates loss  
loss = tf.reduce_mean(tf.square(y - y_train))
```

Minimize loss using an optimizer



`tf.train.GradientDescentOptimizer`

Or, use one of:

`AdadeltaOptimizer`

`AdagradOptimizer`

`AdagradDAOptimizer`

`AdamOptimizer`

...

Automatic Differentiation

Feed (x, y_{label}) pairs and adjust W and b to decrease the loss.

$$W \leftarrow W - \eta (dL/dW)$$

$$b \leftarrow b - \eta (dL/db)$$

TensorFlow computes
gradients automatically

```
# Create an optimizer
```

```
optimizer = tf.train.GradientDescentOptimizer(0.5)
```

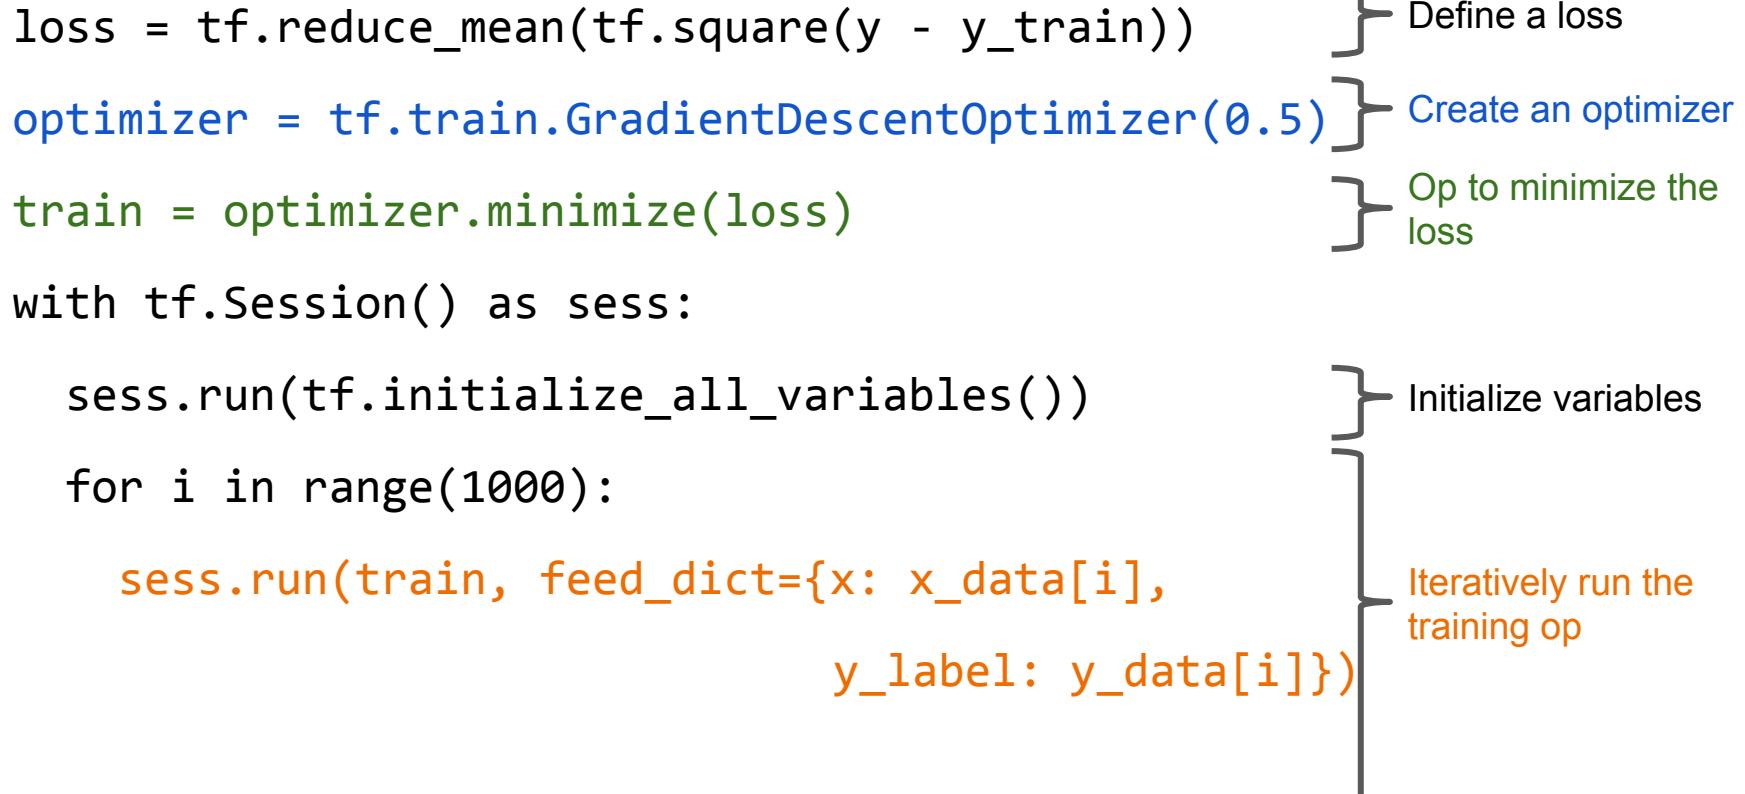
```
# Create an operation that minimizes loss.
```

```
train = optimizer.minimize(loss)
```

Learning rate

Putting it all together

```
loss = tf.reduce_mean(tf.square(y - y_train))  
optimizer = tf.train.GradientDescentOptimizer(0.5)  
train = optimizer.minimize(loss)  
  
with tf.Session() as sess:  
    sess.run(tf.initialize_all_variables())  
  
    for i in range(1000):  
        sess.run(train, feed_dict={x: x_data[i],  
                                  y_label: y_data[i]})
```



The diagram illustrates the components of the provided TensorFlow code. It uses curly braces to group related code snippets:

- A brace on the right side groups the first three lines of code: `loss = tf.reduce_mean(tf.square(y - y_train))`, `optimizer = tf.train.GradientDescentOptimizer(0.5)`, and `train = optimizer.minimize(loss)`. This brace is labeled "Define a loss", "Create an optimizer", and "Op to minimize the loss" respectively.
- A brace on the right side groups the line `sess.run(tf.initialize_all_variables())`. This brace is labeled "Initialize variables".
- A large brace on the right side groups the entire loop structure from `for i in range(1000):` down to the final `sess.run` call. This brace is labeled "Iteratively run the training op".

Questions?

Progress last year

Release

November 2015

December (0.6)

February (0.7)

April (0.8)

June (0.9)

August (0.10)

October (0.11)

November (0.12)

Milestone

Initial release

Faster on GPUs; Python 3.3+

TensorFlow Serving

Distributed TensorFlow

iOS; Mac GPU

Slim

HDFS; CUDA 8, CuDNN 5

Windows 7, 10, and Server 2016

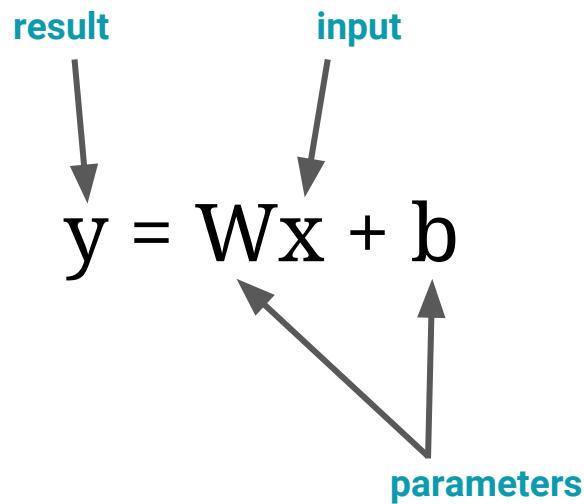
Linear Regression

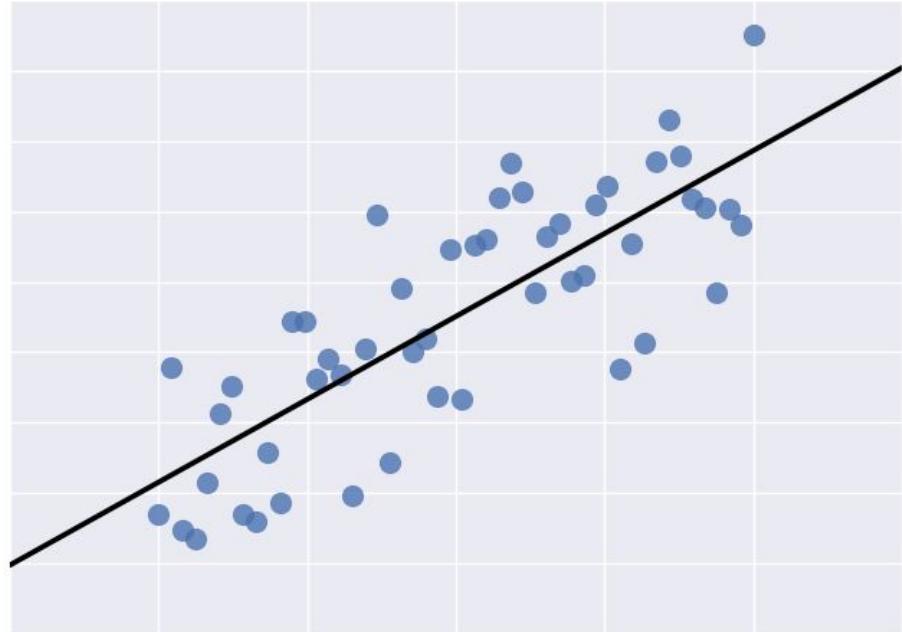
Inference, Loss, Optimize

Linear Regression

$$y = Wx + b$$

result input
parameters





What are we trying to do?

Mystery equation: $y = 0.1 * x + 0.3 + \text{noise}$

Model: $y = W * x + b$

Objective: Given enough (x, y) samples, figure out the value of W and b .

Let's code this up

1. **Inference** (“given x , this is the code to predicts y ”)
2. **Loss** (“quantify how bad our prediction was”)
3. **Optimize** (“update variables to improve the prediction”)

$y = Wx + b$ in TensorFlow

```
import tensorflow as tf
```

$y = Wx + b$ in TensorFlow

```
import tensorflow as tf  
  
x = tf.placeholder(shape=[None],  
                   dtype=tf.float32, name="x")
```

$y = Wx + b$ in TensorFlow

```
import tensorflow as tf

x = tf.placeholder(shape=[None],
                    dtype=tf.float32, name="x")

W = tf.Variable(tf.random_normal([1], name="W")
```

$y = Wx + b$ in TensorFlow

```
import tensorflow as tf

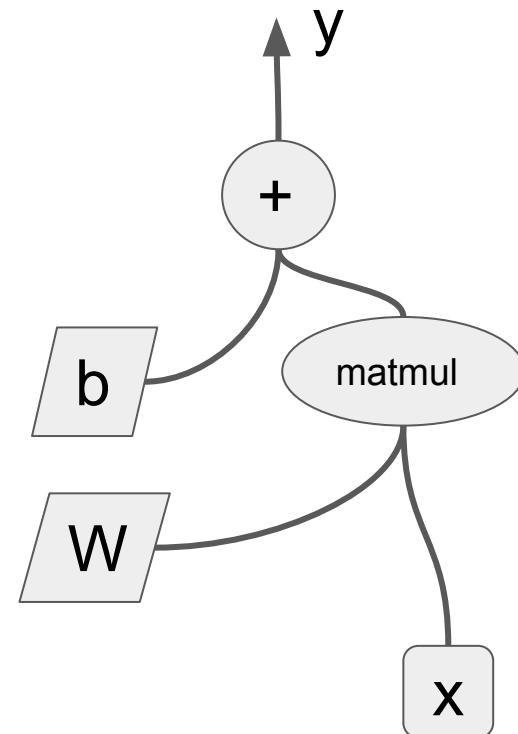
x = tf.placeholder(shape=[None],
                    dtype=tf.float32, name="x")

W = tf.Variable(tf.random_normal([1], name="W"))

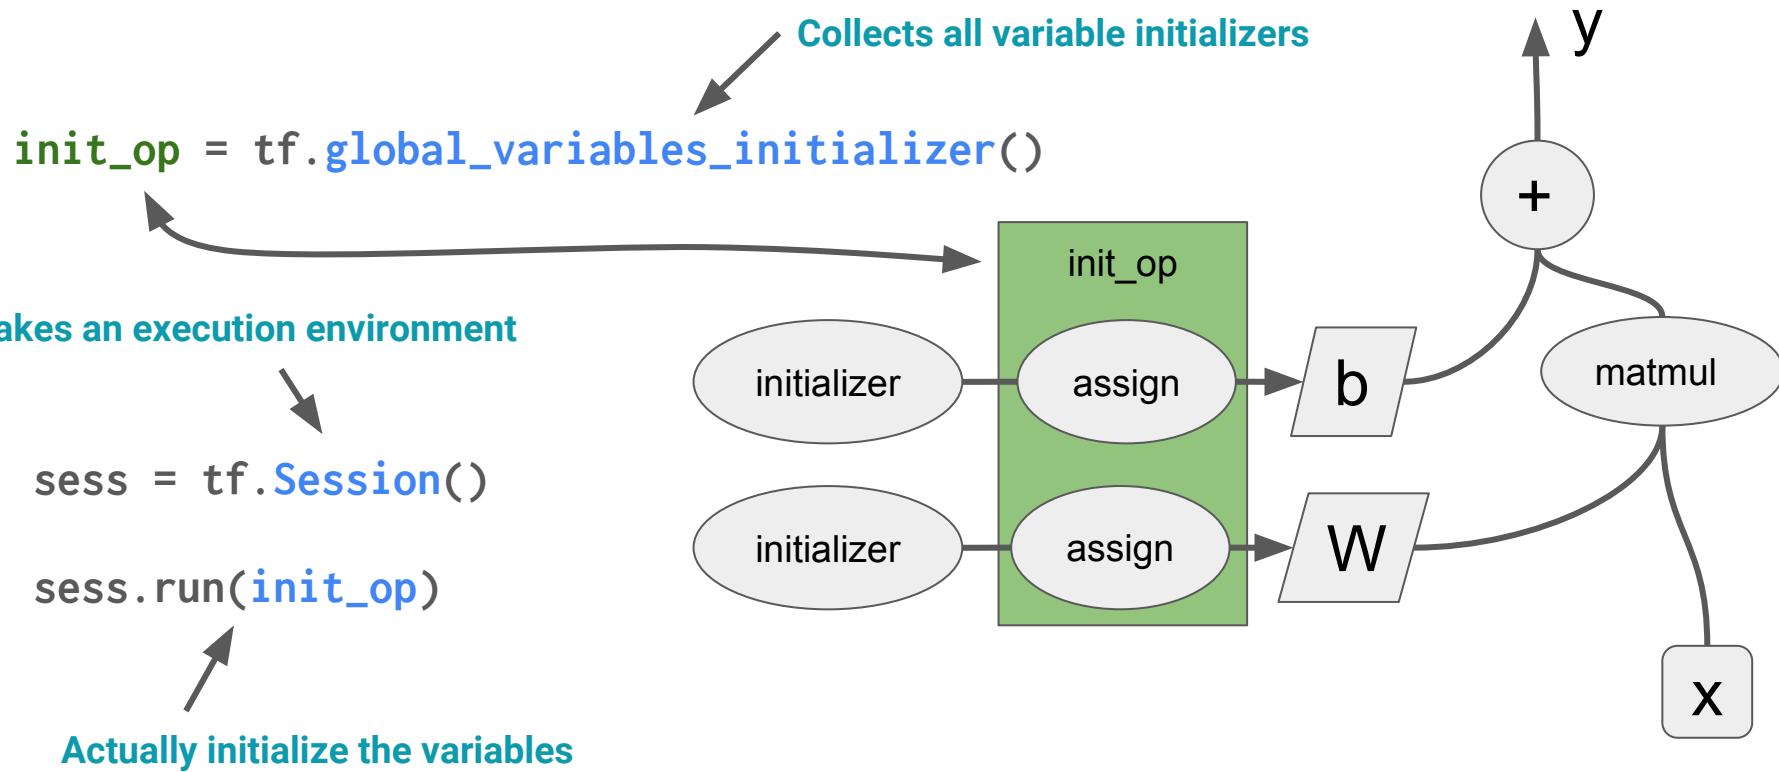
b = tf.Variable(tf.random_normal([1], name="b"))
```

$y = Wx + b$ in TensorFlow

```
import tensorflow as tf  
  
x = tf.placeholder(shape=[None],  
                    dtype=tf.float32, name="x")  
  
W = tf.Variable(tf.random_normal([1], name="W")  
  
b = tf.Variable(tf.random_normal([1], name="b")  
  
y = W * x + b
```



Variables Must be Initialized

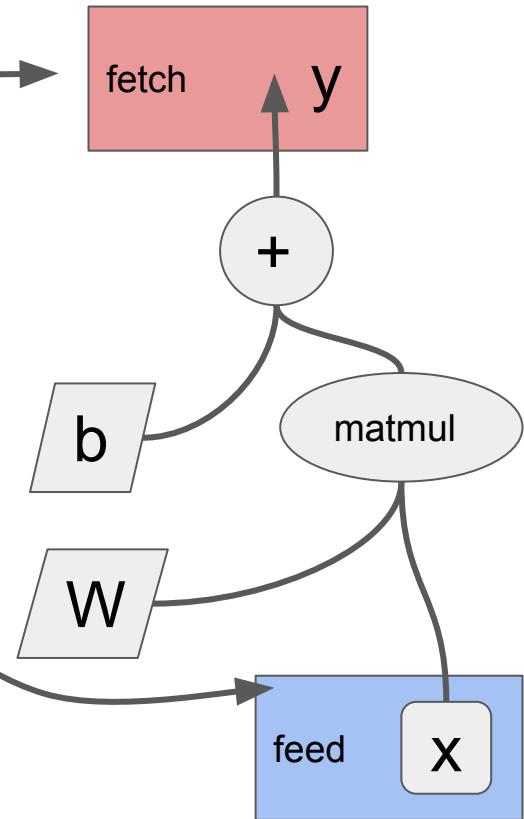


Running the Computation

`x_in = 3`

`sess.run(y, feed_dict={x: x_in})`

- Only what's used to compute a fetch will be evaluated
- All Tensors can be fed, but all placeholders must be fed



Inference: putting it all together

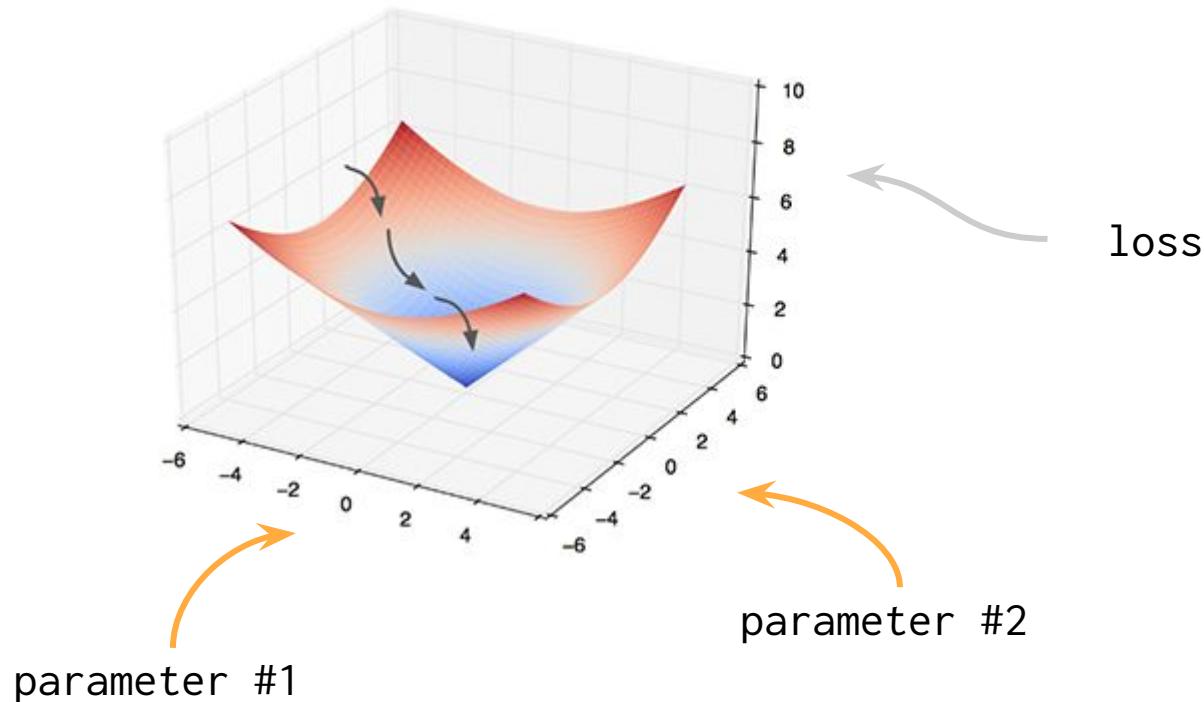
```
import tensorflow as tf
x = tf.placeholder(shape=[None],
                    dtype=tf.float32,
                    name='x')
W = tf.Variable(tf.random_normal([1], name="W"))
b = tf.Variable(tf.random_normal([1], name="b"))
y = W * x + b

with tf.Session() as sess:
    sess.run(tf.initialize_all_variables())
    print(sess.run(y, feed_dict={x: x_in}))
```

The diagram illustrates the execution flow of the provided TensorFlow code. It is divided into four main phases, each indicated by a curly brace on the right side:

- Build the graph**: This phase covers the declaration of tensors and variables. It includes the definition of `x` as a placeholder, the creation of `W` and `b` as variables, and the computation of `y` as the product of `W` and `x` plus `b`.
- Prepare execution env**: This phase covers the setup of the execution environment. It includes the creation of a `Session` object.
- Initialize variables**: This phase covers the initialization of all variables. It includes the execution of the `tf.initialize_all_variables()` operation.
- Run the computation**: This phase covers the final computation. It includes the execution of the `sess.run(y, feed_dict={x: x_in})` operation.

Concepts: Loss and Gradient Descent



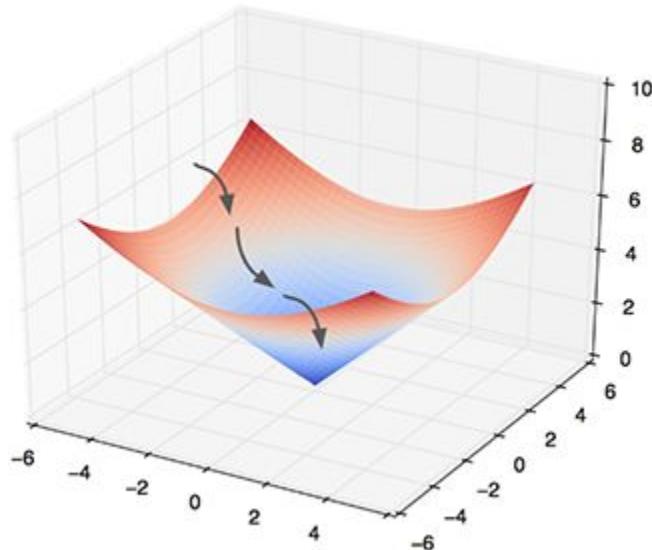
Define a Loss

Given y , y_{train} compute a loss, for instance:

$$loss = (y - y_{train})^2$$

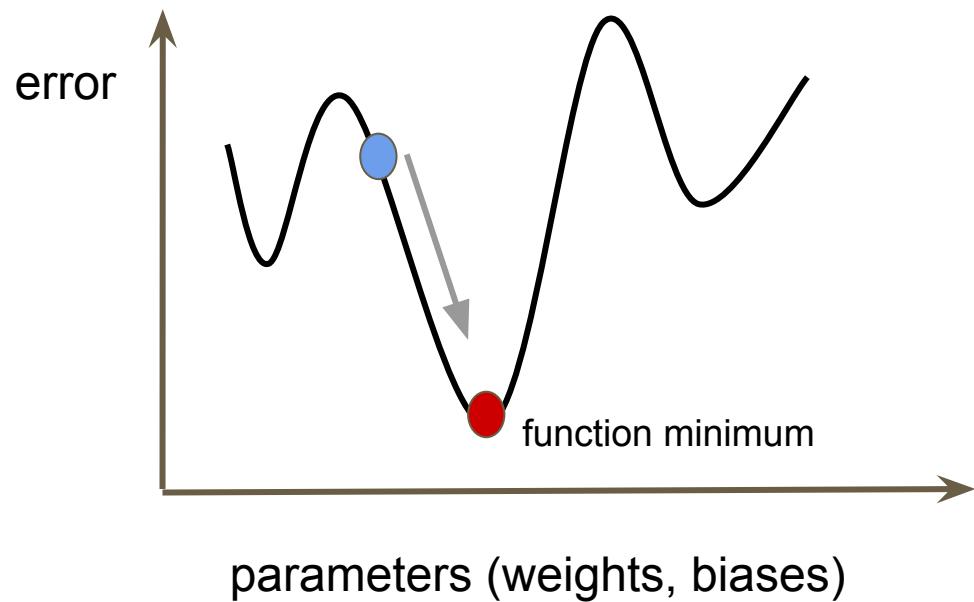
```
# create an operation that calculates loss  
loss = tf.reduce_mean(tf.square(y - y_train))
```

Optimization



-5	-4	-3	-2	-1	0	1	2	3	4	5
5	50	41	34	29	26	25	26	29	34	41
4	41	32	25	20	17	16	17	20	25	32
3	34	25	18	13	10	9	10	13	18	25
2	29	20	13	8	5	4	5	8	13	20
1	26	17	10	5	2	1	2	5	10	17
0	25	16	9		1	0	1	4	9	16
-1	26	17	10	5	2	1	2	5	10	17
-2	29	20	13	8	5	4	5	8	13	20
-3	34	25	18	13	10	9	10	13	18	25
-4	41	32	25	20	17	16	17	20	25	32
-5	50	41	34	29	26	25	26	29	34	41

Minimize loss using an optimizer



`tf.train.GradientDescentOptimizer`

Or, use one of:

`AdadeltaOptimizer`

`AdagradOptimizer`

`AdagradDAOptimizer`

`AdamOptimizer`

...

Automatic Differentiation

Feed (x, y_{label}) pairs and adjust W and b to decrease the loss.

$$W \leftarrow W - \eta (dL/dW)$$

$$b \leftarrow b - \eta (dL/db)$$

TensorFlow computes
gradients automatically

```
# Create an optimizer
```

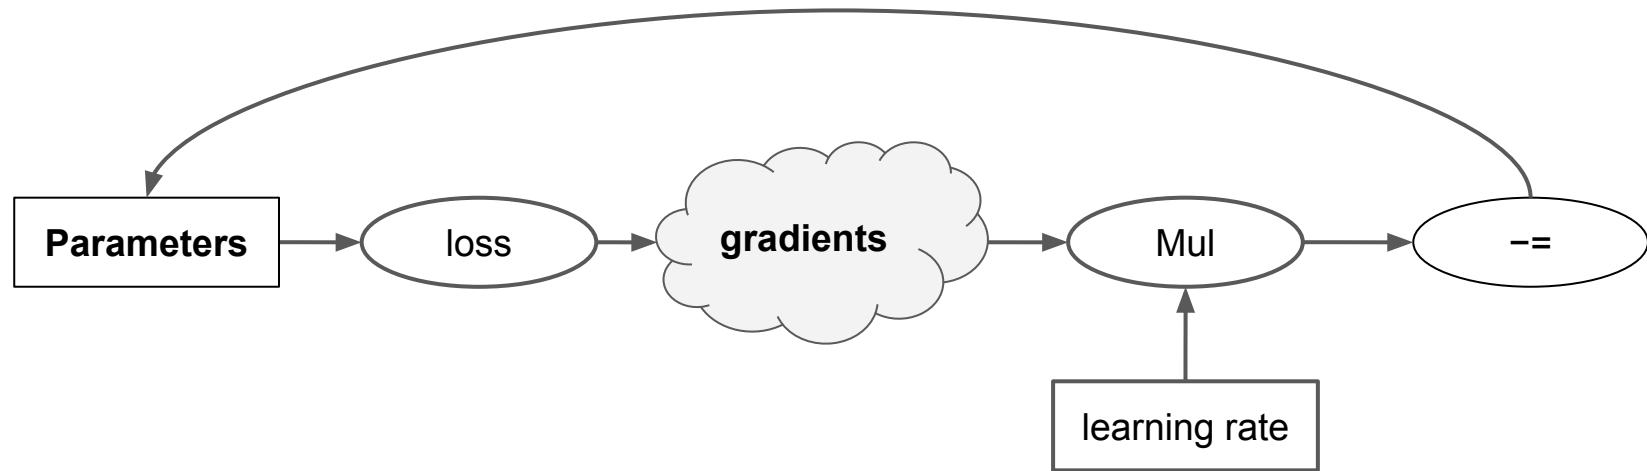
```
optimizer = tf.train.GradientDescentOptimizer(0.5)
```

```
# Create an operation that minimizes loss.
```

```
train = optimizer.minimize(loss)
```

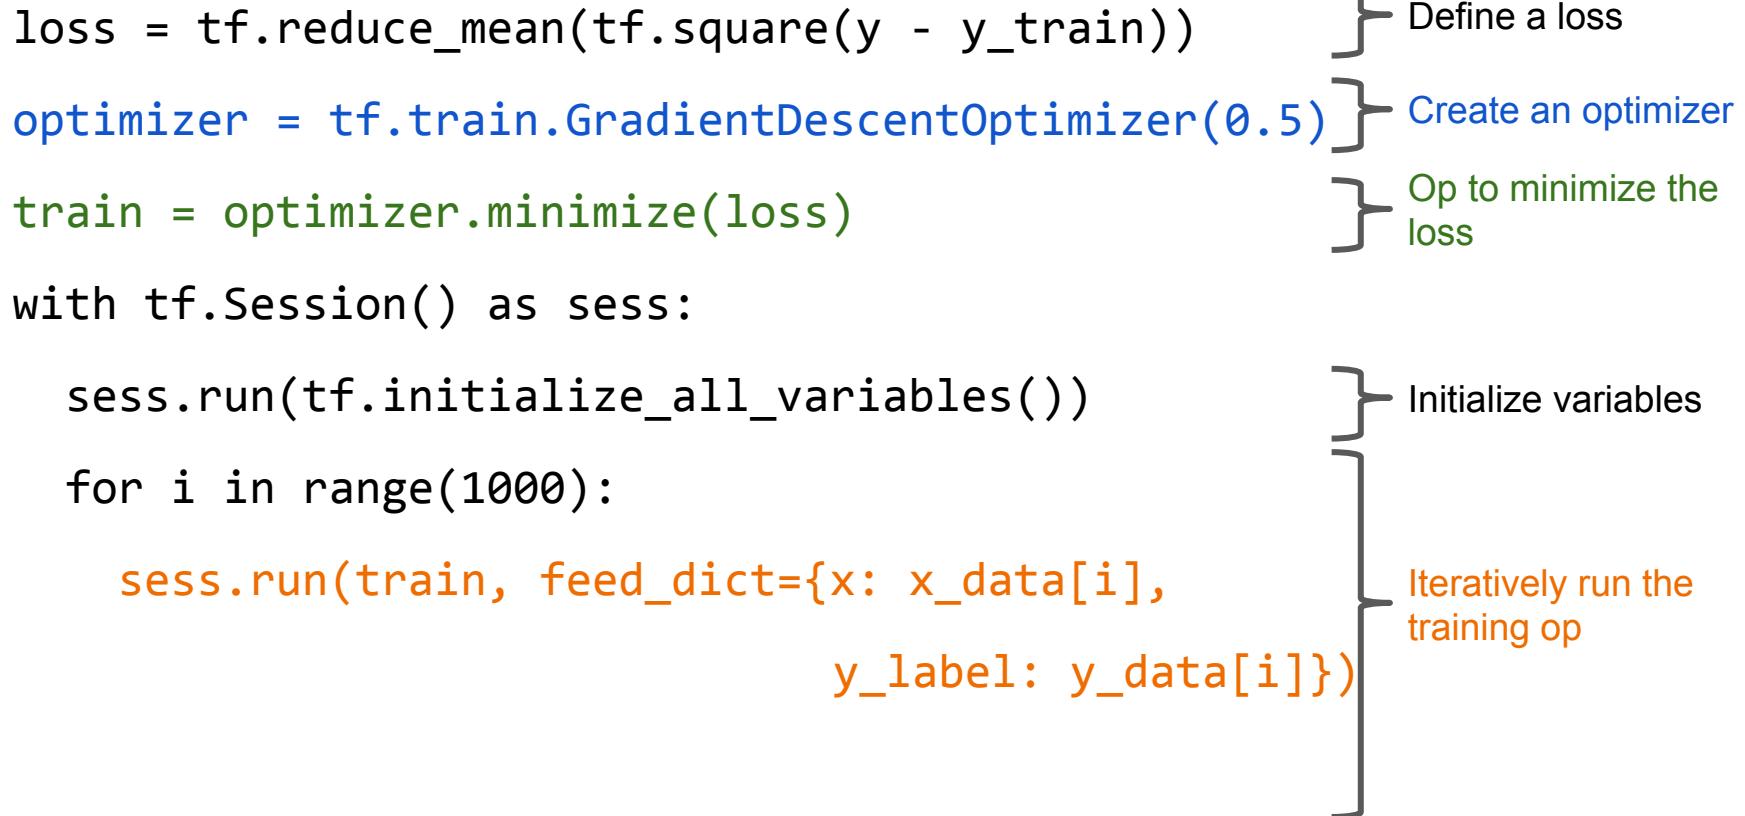
Learning rate

Behind the scenes



Putting it all together

```
loss = tf.reduce_mean(tf.square(y - y_train))  
optimizer = tf.train.GradientDescentOptimizer(0.5)  
train = optimizer.minimize(loss)  
  
with tf.Session() as sess:  
    sess.run(tf.initialize_all_variables())  
  
    for i in range(1000):  
        sess.run(train, feed_dict={x: x_data[i],  
                                  y_label: y_data[i]})
```



The diagram illustrates the components of the provided TensorFlow code. It uses curly braces to group related code snippets:

- A brace on the right side groups the first three lines of code: `loss = tf.reduce_mean(tf.square(y - y_train))`, `optimizer = tf.train.GradientDescentOptimizer(0.5)`, and `train = optimizer.minimize(loss)`. This brace is labeled "Define a loss", "Create an optimizer", and "Op to minimize the loss" respectively.
- A brace on the right side groups the line `sess.run(tf.initialize_all_variables())`. This brace is labeled "Initialize variables".
- A large brace on the right side groups the entire loop structure from `for i in range(1000):` down to the final `sess.run` call. This brace is labeled "Iteratively run the training op".

Cumulative GitHub Stars

